

lab12

October 6, 2023

0.1 Lab 12: Unsupervised Learning

Jack Krebsbach Math 313

0.1.1 12.5.1 Principal Components Analysis

First we explore the USArrests data set.

```
[ ]: ## All of our imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.datasets import get_rdataset
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from ISLP import load_data
```

```
[ ]: ## New imports needed
from sklearn.cluster import (KMeans, AgglomerativeClustering)
from scipy.cluster.hierarchy import (dendrogram, cut_tree)
from ISLP.cluster import compute_linkage
```

There are 50 rows in this data set, which contain the states in alphabetical order.

```
[ ]: USArrests = get_rdataset('USArrests').data
USArrests
```

```
[ ]:
```

	Murder	Assault	UrbanPop	Rape
rownames				
Alabama	13.2	236	58	21.2
Alaska	10.0	263	48	44.5
Arizona	8.1	294	80	31.0
Arkansas	8.8	190	50	19.5
California	9.0	276	91	40.6
Colorado	7.9	204	78	38.7
Connecticut	3.3	110	77	11.1
Delaware	5.9	238	72	15.8
Florida	15.4	335	80	31.9
Georgia	17.4	211	60	25.8

Hawaii	5.3	46	83	20.2
Idaho	2.6	120	54	14.2
Illinois	10.4	249	83	24.0
Indiana	7.2	113	65	21.0
Iowa	2.2	56	57	11.3
Kansas	6.0	115	66	18.0
Kentucky	9.7	109	52	16.3
Louisiana	15.4	249	66	22.2
Maine	2.1	83	51	7.8
Maryland	11.3	300	67	27.8
Massachusetts	4.4	149	85	16.3
Michigan	12.1	255	74	35.1
Minnesota	2.7	72	66	14.9
Mississippi	16.1	259	44	17.1
Missouri	9.0	178	70	28.2
Montana	6.0	109	53	16.4
Nebraska	4.3	102	62	16.5
Nevada	12.2	252	81	46.0
New Hampshire	2.1	57	56	9.5
New Jersey	7.4	159	89	18.8
New Mexico	11.4	285	70	32.1
New York	11.1	254	86	26.1
North Carolina	13.0	337	45	16.1
North Dakota	0.8	45	44	7.3
Ohio	7.3	120	75	21.4
Oklahoma	6.6	151	68	20.0
Oregon	4.9	159	67	29.3
Pennsylvania	6.3	106	72	14.9
Rhode Island	3.4	174	87	8.3
South Carolina	14.4	279	48	22.5
South Dakota	3.8	86	45	12.8
Tennessee	13.2	188	59	26.9
Texas	12.7	201	80	25.5
Utah	3.2	120	80	22.9
Vermont	2.2	48	32	11.2
Virginia	8.5	156	63	20.7
Washington	4.0	145	73	26.2
West Virginia	5.7	81	39	9.3
Wisconsin	2.6	53	66	10.8
Wyoming	6.8	161	60	15.6

```
[ ]: USArrests.columns
```

```
[ ]: Index(['Murder', 'Assault', 'UrbanPop', 'Rape'], dtype='object')
```

```
[ ]: # The variables have lots of different means.
      USArrests.mean()
```

```
[ ]: Murder      7.788
      Assault    170.760
      UrbanPop   65.540
      Rape       21.232
      dtype: float64
```

```
[ ]: # We can compute statistics like variance on these features
      USArrests.var()
```

```
[ ]: Murder      18.970465
      Assault    6945.165714
      UrbanPop   209.518776
      Rape       87.729159
      dtype: float64
```

```
[ ]: # We can standardize the data.
      scaler = StandardScaler(with_std=True, with_mean=True)
      USArrests_scaled = scaler.fit_transform(USArrests)
      USArrests_scaled
```

```
[ ]: array([[ 1.25517927,  0.79078716, -0.52619514, -0.00345116],
            [ 0.51301858,  1.11805959, -1.22406668,  2.50942392],
            [ 0.07236067,  1.49381682,  1.00912225,  1.05346626],
            [ 0.23470832,  0.23321191, -1.08449238, -0.18679398],
            [ 0.28109336,  1.2756352 ,  1.77678094,  2.08881393],
            [ 0.02597562,  0.40290872,  0.86954794,  1.88390137],
            [-1.04088037, -0.73648418,  0.79976079, -1.09272319],
            [-0.43787481,  0.81502956,  0.45082502, -0.58583422],
            [ 1.76541475,  1.99078607,  1.00912225,  1.1505301 ],
            [ 2.22926518,  0.48775713, -0.38662083,  0.49265293],
            [-0.57702994, -1.51224105,  1.21848371, -0.11129987],
            [-1.20322802, -0.61527217, -0.80534376, -0.75839217],
            [ 0.60578867,  0.94836277,  1.21848371,  0.29852525],
            [-0.13637203, -0.70012057, -0.03768506, -0.0250209 ],
            [-1.29599811, -1.39102904, -0.5959823 , -1.07115345],
            [-0.41468229, -0.67587817,  0.03210209, -0.34856705],
            [ 0.44344101, -0.74860538, -0.94491807, -0.53190987],
            [ 1.76541475,  0.94836277,  0.03210209,  0.10439756],
            [-1.31919063, -1.06375661, -1.01470522, -1.44862395],
            [ 0.81452136,  1.56654403,  0.10188925,  0.70835037],
            [-0.78576263, -0.26375734,  1.35805802, -0.53190987],
            [ 1.00006153,  1.02108998,  0.59039932,  1.49564599],
            [-1.1800355 , -1.19708982,  0.03210209, -0.68289807],
            [ 1.9277624 ,  1.06957478, -1.5032153 , -0.44563089],
            [ 0.28109336,  0.0877575 ,  0.31125071,  0.75148985],
            [-0.41468229, -0.74860538, -0.87513091, -0.521125 ],
            [-0.80895515, -0.83345379, -0.24704653, -0.51034012],
```

```
[ 1.02325405,  0.98472638,  1.0789094 ,  2.671197  ],
[-1.31919063, -1.37890783, -0.66576945, -1.26528114],
[-0.08998698, -0.14254532,  1.63720664, -0.26228808],
[ 0.83771388,  1.38472601,  0.31125071,  1.17209984],
[ 0.76813632,  1.00896878,  1.42784517,  0.52500755],
[ 1.20879423,  2.01502847, -1.43342815, -0.55347961],
[-1.62069341, -1.52436225, -1.5032153 , -1.50254831],
[-0.11317951, -0.61527217,  0.66018648,  0.01811858],
[-0.27552716, -0.23951493,  0.1716764 , -0.13286962],
[-0.66980002, -0.14254532,  0.10188925,  0.87012344],
[-0.34510472, -0.78496898,  0.45082502, -0.68289807],
[-1.01768785,  0.03927269,  1.49763233, -1.39469959],
[ 1.53348953,  1.3119988 , -1.22406668,  0.13675217],
[-0.92491776, -1.027393 , -1.43342815, -0.90938037],
[ 1.25517927,  0.20896951, -0.45640799,  0.61128652],
[ 1.13921666,  0.36654512,  1.00912225,  0.46029832],
[-1.06407289, -0.61527217,  1.00912225,  0.17989166],
[-1.29599811, -1.48799864, -2.34066115, -1.08193832],
[ 0.16513075, -0.17890893, -0.17725937, -0.05737552],
[-0.87853272, -0.31224214,  0.52061217,  0.53579242],
[-0.48425985, -1.08799901, -1.85215107, -1.28685088],
[-1.20322802, -1.42739264,  0.03210209, -1.1250778 ],
[-0.22914211, -0.11830292, -0.38662083, -0.60740397]]])
```

Once we scale the data we can now perform PCA on the data.

```
[ ]: np.std(USArrests_scaled[:, 1])

[ ]: 1.0

[ ]: ## This uses the PCA package versus by hand
pcaUS = PCA()

[ ]: pcaUS.fit(USArrests_scaled)

[ ]: PCA()

[ ]: pcaUS.mean_

[ ]: array([-7.10542736e-17,  1.38777878e-16, -4.39648318e-16,  8.59312621e-16])

[ ]: scores = pcaUS.transform(USArrests_scaled)
```

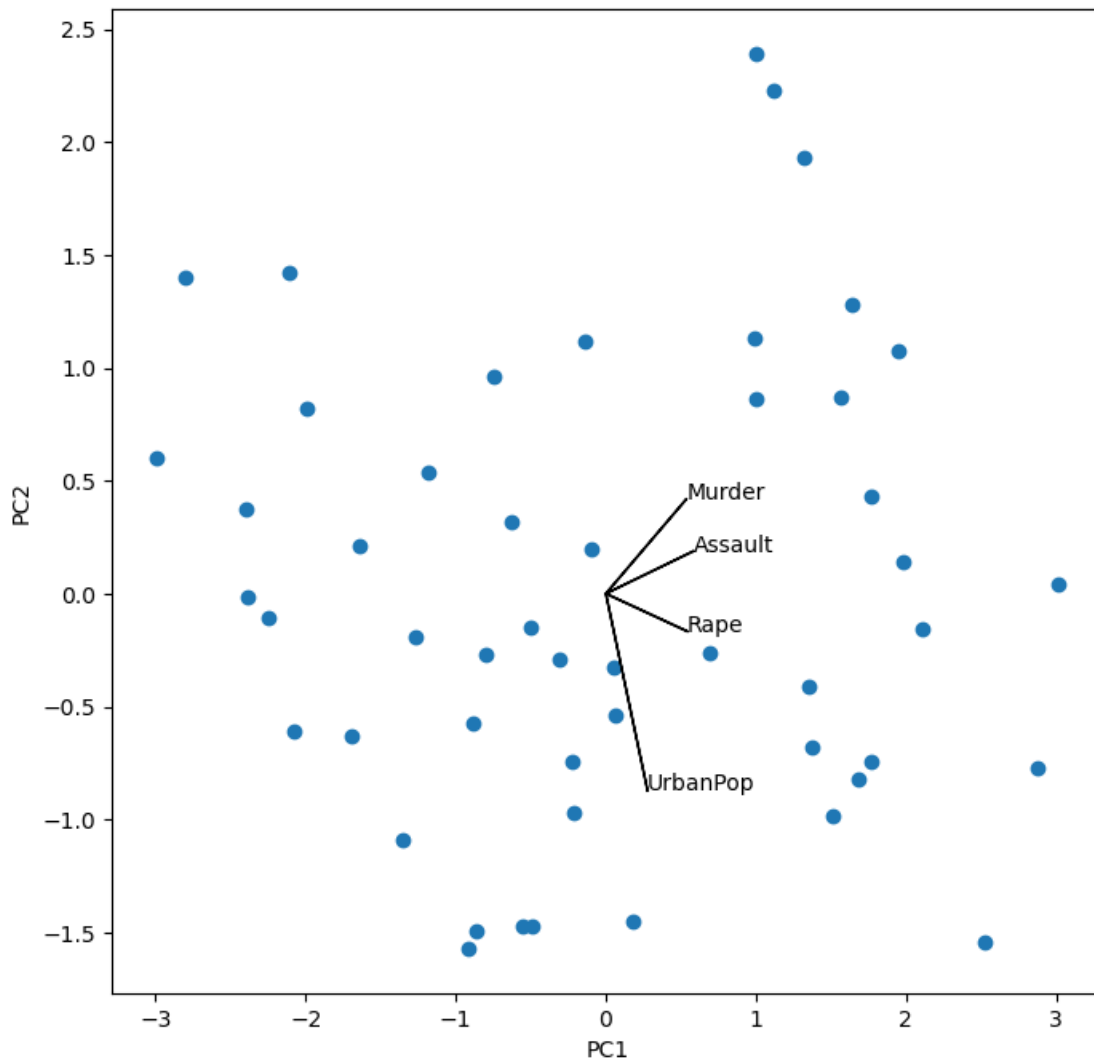
Principal Component Loadings The loadings of the principal component analysis can be extracted using `.components_` of the PCA analysis. Each row contains the PC loading vector.

```
[ ]: pcaUS.components_
```

```
[ ]: array([[ 0.53589947,  0.58318363,  0.27819087,  0.54343209],
          [ 0.41818087,  0.1879856 , -0.87280619, -0.16731864],
          [-0.34123273, -0.26814843, -0.37801579,  0.81777791],
          [ 0.6492278 , -0.74340748,  0.13387773,  0.08902432]])
```

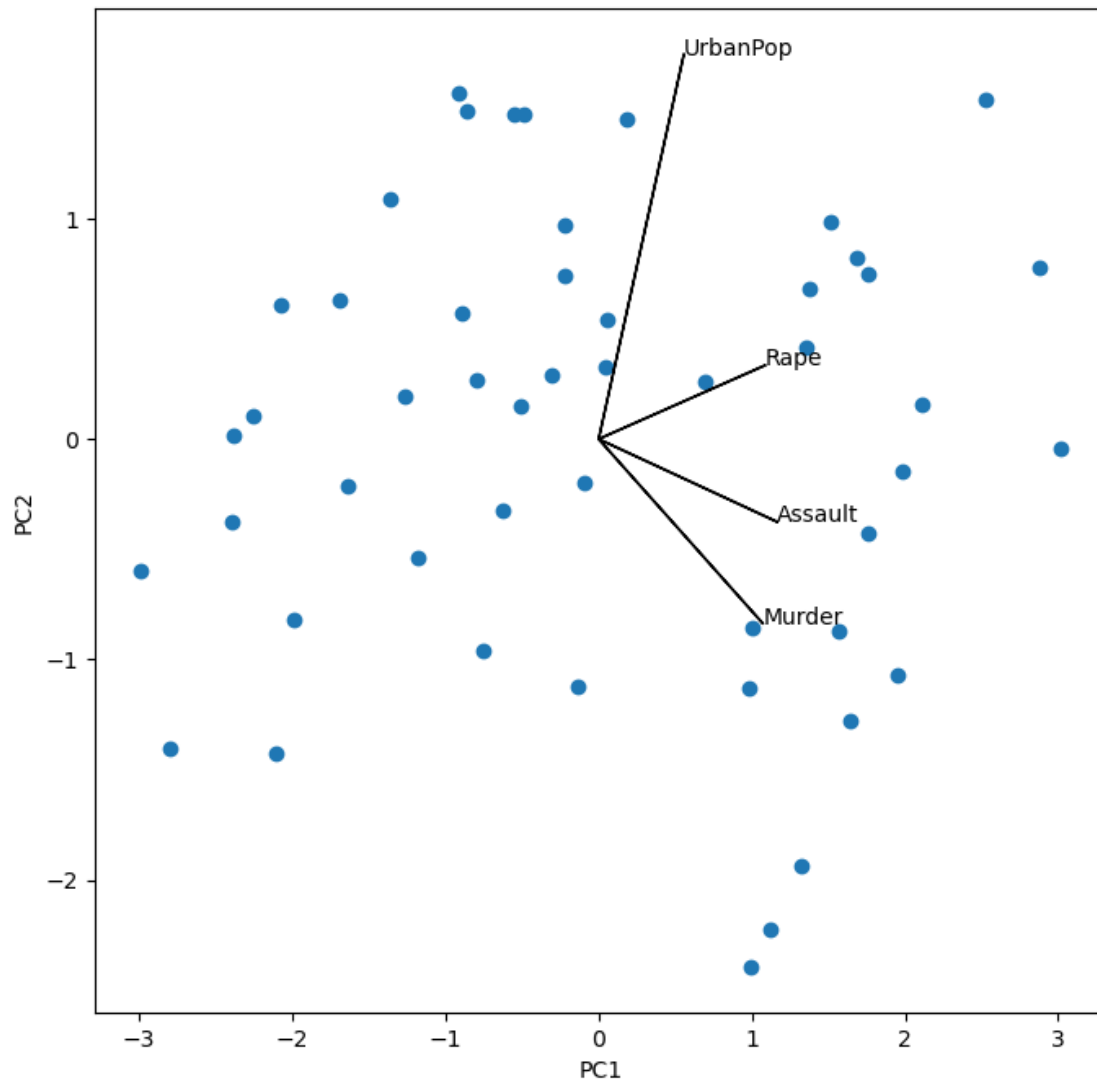
```
[ ]: # Here we make the bi-plot manually as it is not a standard import from sklearn.
```

```
i, j = 0, 1 # which components
fig, ax = plt.subplots(1, 1, figsize=(8, 8))
ax.scatter(scores[:,0], scores[:,1])
ax.set_xlabel('PC%d' % (i+1))
ax.set_ylabel('PC%d' % (j+1))
for k in range(pcaUS.components_.shape[1]):
    ax.arrow(0, 0, pcaUS.components_[i,k], pcaUS.components_[j,k])
    ax.text(pcaUS.components_[i,k], pcaUS.components_[j,k], USArrests.
columns[k])
```



This figure is a reflection (through the *y*-axis) from the book in figure 12.1. The PC are unique only up to sign change. Thus, we can replicate the figure by flipping the signs and can more clearly see the differences between loadings by increasing the length of the arrows.

```
[ ]: scale_arrow = s_ = 2
scores[:,1] *= -1
pcaUS.components_[1] *= -1 # flip the y-axis
fig, ax = plt.subplots(1, 1, figsize=(8, 8))
ax.scatter(scores[:,0], scores[:,1])
ax.set_xlabel('PC%d' % (i+1))
ax.set_ylabel('PC%d' % (j+1))
for k in range(pcaUS.components_.shape[1]):
    ax.arrow(0, 0, s_*pcaUS.components_[i,k], s_*pcaUS.components_[ j,k])
    ax.text(s_*pcaUS.components_[i,k], s_*pcaUS.components_[j,k],
           USArrests.columns[k])
```



```
[ ]: # We grab the standard deviations
scores.std(0, ddof=1)
```

```
[ ]: array([1.5908673 , 1.00496987, 0.6031915 , 0.4206774 ])
```

What we are really after is how much variance each PC can explain, here we can find the explained variance by component. It is helpful to also find the ratio of explained variance to get a better of idea this.

```
[ ]: pcaUS.explained_variance_
```

```
[ ]: array([2.53085875, 1.00996444, 0.36383998, 0.17696948])
```

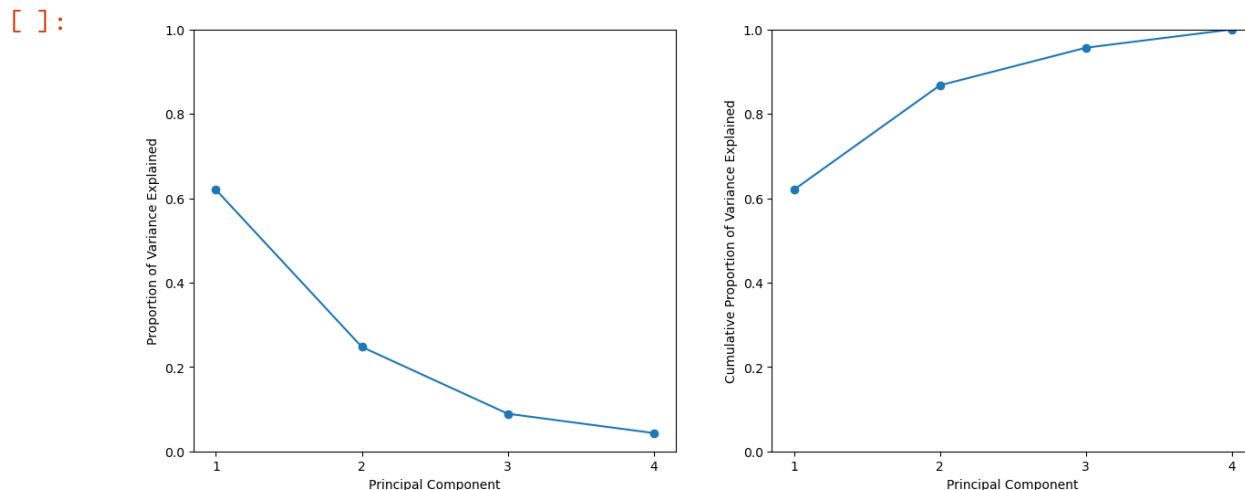
```
[ ]: pcaUS.explained_variance_ratio_
```

```
[ ]: array([0.62006039, 0.24744129, 0.0891408 , 0.04335752])
```

The first principal component explains 60% of the variance in the data. To visually see this we can plot it below.

```
[ ]: %%capture
fig, axes = plt.subplots(1, 2, figsize=(15, 6))
ticks = np.arange(pcaUS.n_components_)+1
ax = axes[0]
ax.plot(ticks,
pcaUS.explained_variance_ratio_ , marker='o')
ax.set_xlabel('Principal Component'); ax.set_ylabel('Proportion of Variance_
↳Explained')
ax.set_ylim([0,1])
ax.set_xticks(ticks)
ax.set_xlabel('Principal Component'); ax.set_ylabel('Proportion of Variance_
↳Explained')
ax.set_ylim([0,1])
ax.set_xticks(ticks)
```

```
[ ]: ax = axes[1]
ax.plot(ticks,
pcaUS.explained_variance_ratio_.cumsum(),
marker='o')
ax.set_xlabel('Principal Component')
ax.set_ylabel('Cumulative Proportion of Variance Explained')
ax.set_ylim([0, 1])
ax.set_xticks(ticks)
fig
```




```
[ ]: a = np.array([1,2,8,-3])
      np.cumsum(a)
```

```
[ ]: array([ 1,  3, 11,  8])
```

0.1.2 12.5.2 Matrix Completion

Our goal in this section is to re-create the analysis of the data set in section 12.3. Here we use Singular Value Decomposition to solve for the principal components of the data.

```
[ ]: X = USArrests_scaled
      ## These are the matrix that our matrix A can be decomposed into.
      U, D, V = np.linalg.svd(X, full_matrices=False)
      U.shape, D.shape, V.shape
```

```
[ ]: ((50, 4), (4,), (4, 4))
```

```
[ ]: #V is equivalent to the loading matrix from the PCA before!
      V
```

```
[ ]: array([[ -0.53589947, -0.58318363, -0.27819087, -0.54343209],
            [ -0.41818087, -0.1879856 ,  0.87280619,  0.16731864],
            [  0.34123273,  0.26814843,  0.37801579, -0.81777791],
            [  0.6492278 , -0.74340748,  0.13387773,  0.08902432]])
```

```
[ ]: pcaUS.components_
```

```
[ ]: array([[ 0.53589947,  0.58318363,  0.27819087,  0.54343209],
            [-0.41818087, -0.1879856 ,  0.87280619,  0.16731864],
            [-0.34123273, -0.26814843, -0.37801579,  0.81777791],
            [ 0.6492278 , -0.74340748,  0.13387773,  0.08902432]])
```

The matrix U is a standardized version of the PCA score matrix. The standardization involves scaling the columns to have sum-of-squares 1.

```
[ ]: (U * D[None,:])[:3]
```

```
[ ]: array([[ -0.98556588, -1.13339238,  0.44426879,  0.15626714],
            [-1.95013775, -1.07321326, -2.04000333, -0.43858344],
            [-1.76316354,  0.74595678, -0.05478082, -0.83465292]])
```

```
[ ]: scores[:3]
```

```
[ ]: array([[ 0.98556588, -1.13339238, -0.44426879,  0.15626714],
            [ 1.95013775, -1.07321326,  2.04000333, -0.43858344],
            [ 1.76316354,  0.74595678,  0.05478082, -0.83465292]])
```

Note: this section of the lab would be able to be done by just using the `PCA()` estimator but we are interested in using `np.linalg.svd()` to help us explore how matrix completion works.

We purposely omit 50 random entries in our matrix before implementing the algorithm from 12.1 for matrix completion.

```
[ ]: n_omit = 20
      np.random.seed(15)
      r_idx = np.random.choice(np.arange(X.shape[0]), n_omit ,replace=False)
      c_idx = np.random.choice(np.arange(X.shape[1]), n_omit , replace=True)
      Xna = X.copy()
      Xna[r_idx, c_idx] = np.nan
```

This function allows us to take in a matrix and return its estimation from SVD.

```
[ ]: def low_rank(X, M=1):
      U, D, V = np.linalg.svd(X)
      L = U[:, :M] * D[None, :M]
      return L.dot(V[:M])
```

Here we replace the missing values (the ones we purposely got rid of) with the means of the other entries in the column.

The matrix `imiss` is a logical matrix having the same dimensions of `Xna`.

Step 2A involved approximating `Xhat` with our function defined `low_rank()`. Set 2B we use `Xapp` to update the estimates for the missing data. In step 3C we finally compute the relative error.

```
[ ]: Xhat = Xna.copy()
      Xbar = np.nanmean(Xhat, axis=0)
      Xhat[r_idx, c_idx] = Xbar[c_idx]
```

```
[ ]: thresh = 1e-7
      rel_err = 1
      count = 0
      ismiss = np.isnan(Xna)
      mssold = np.mean(Xhat[~issmiss]**2)
      mss0 = np.mean(Xna[~issmiss]**2)
```

This process actually only takes eight iterations as after that the relative error falls below `1e-7`.

```
[ ]: while rel_err > thresh:
      count += 1
      # Step 2(a)
      Xapp = low_rank(Xhat, M=1)
      # Step 2(b)
      Xhat[issmiss] = Xapp[issmiss]
      # Step 2(c)
      mss = np.mean(((Xna - Xapp)[~issmiss])**2)
      rel_err = (mssold - mss) / mss0
```

```

mssold = mss
print("Iteration: {0}, MSS:{1:.3f}, Rel.Err {2:.2e}".format(count, mss,
↪rel_err))

```

```

Iteration: 1, MSS:0.395, Rel.Err 5.99e-01
Iteration: 2, MSS:0.382, Rel.Err 1.33e-02
Iteration: 3, MSS:0.381, Rel.Err 1.44e-03
Iteration: 4, MSS:0.381, Rel.Err 1.79e-04
Iteration: 5, MSS:0.381, Rel.Err 2.58e-05
Iteration: 6, MSS:0.381, Rel.Err 4.22e-06
Iteration: 7, MSS:0.381, Rel.Err 7.65e-07
Iteration: 8, MSS:0.381, Rel.Err 1.48e-07
Iteration: 9, MSS:0.381, Rel.Err 2.95e-08

```

```
[ ]: np.corrcoef(Xapp[ismiss], X[ismiss])[0,1]
```

```
[ ]: 0.7113567434297362
```

0.1.3 12.5.3 Clustering

K-Means Clustering We can use the estimator from the `sklearn` package to perform K-means clustering.

```
[ ]: np.random.seed(0);
X = np.random.standard_normal((50,2)); X[:25,0] += 3;
X[:25,1] -= 4;
```

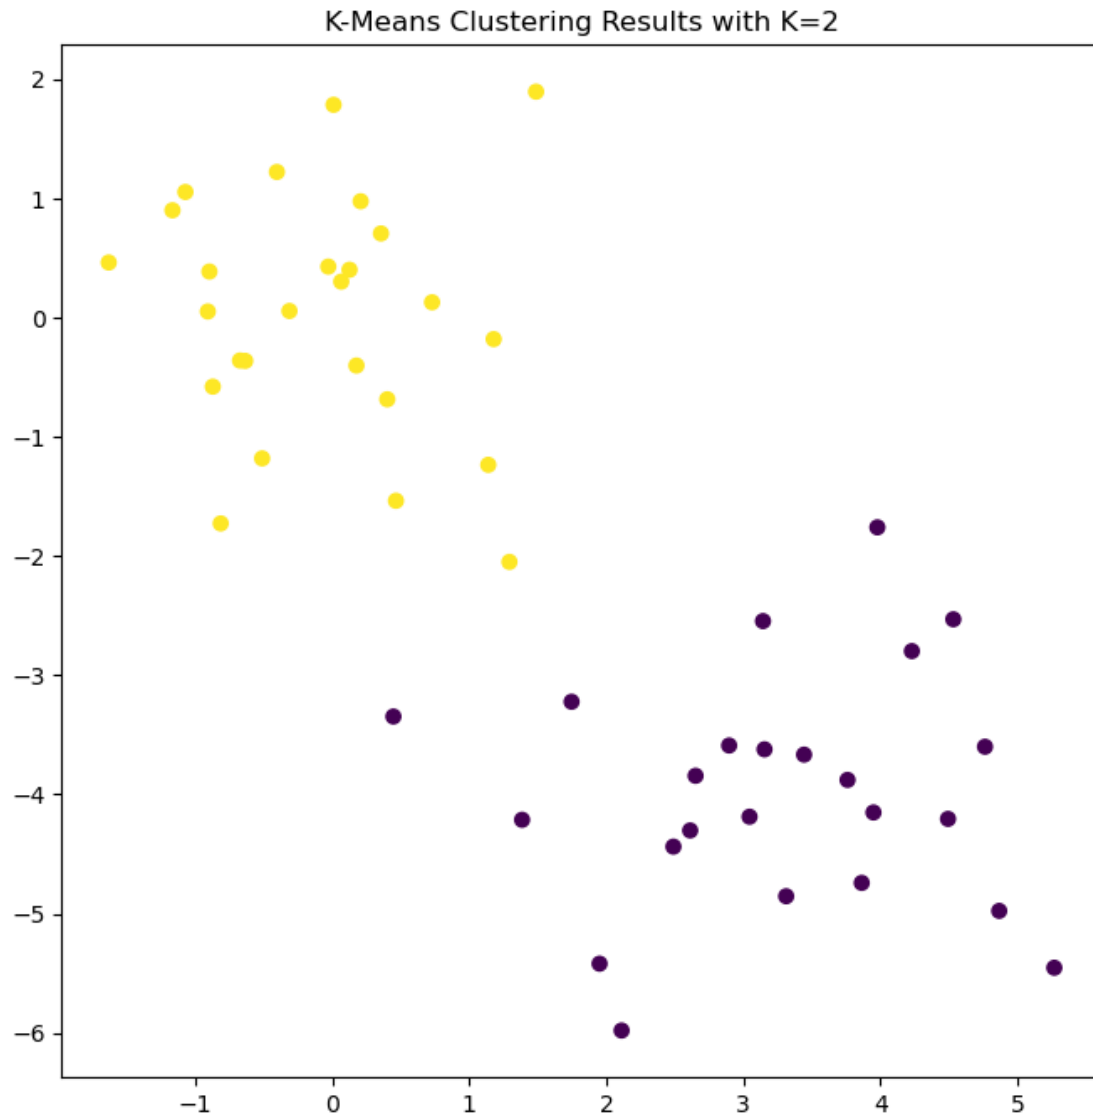
```
[ ]: # This performs k-means with 2 clusters
kmeans = KMeans(n_clusters=2, random_state=2, n_init=20).fit(X)
# Kmeans with 5 clusters
kmeans_three_clsuters = KMeans(n_clusters=5, random_state=2, n_init=20).fit(X)
```

```
[ ]: # It looks like we were able to split teh data perfectly even though we did not
↪have any labels!
kmeans.labels_
```

```
[ ]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
          0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
          1, 1, 1, 1, 1, 1], dtype=int32)
```

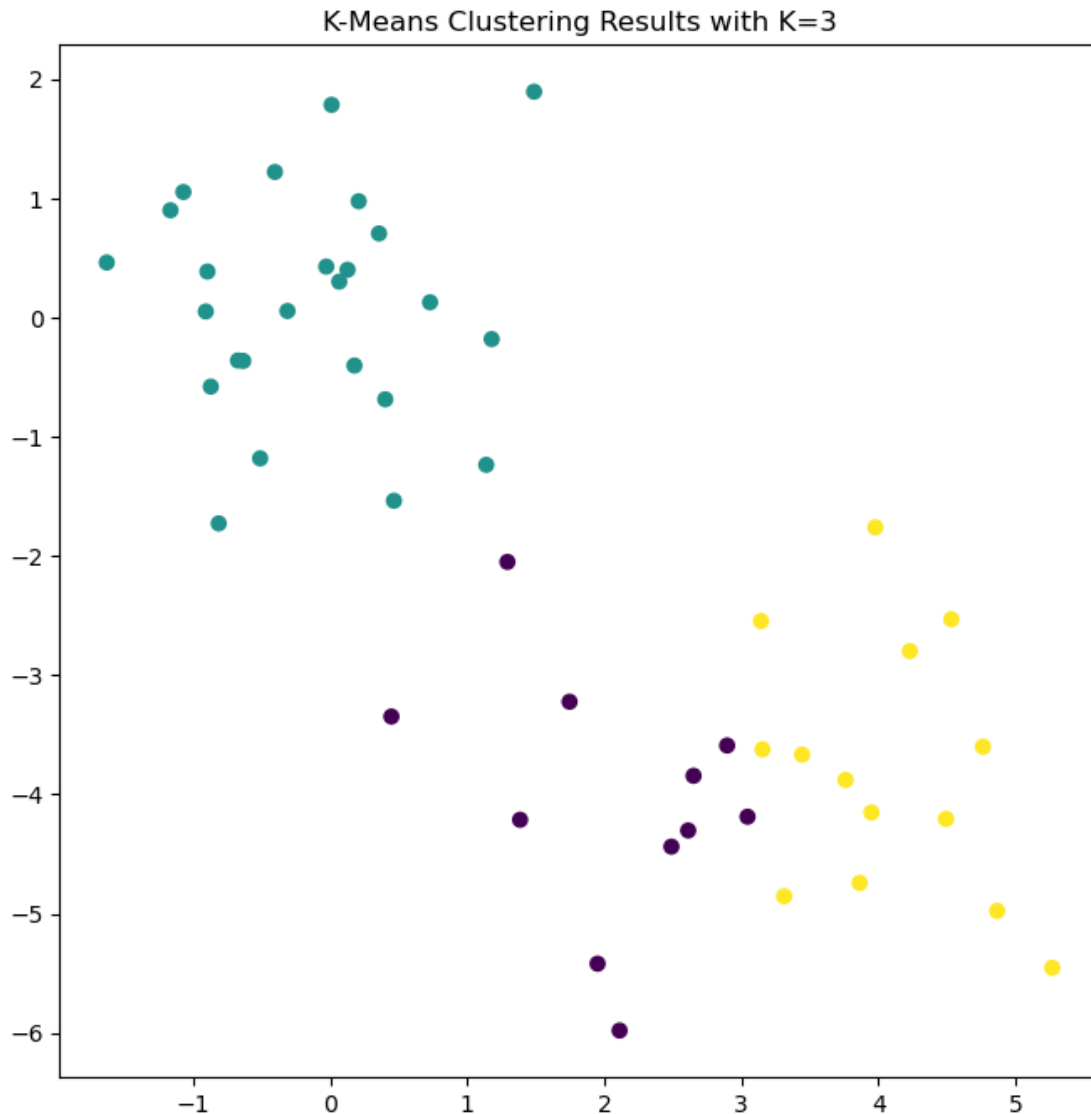
It is easy to plot the observations in the following cell because they are already two-dimensional.

```
[ ]: fig, ax = plt.subplots(1, 1, figsize=(8,8))
ax.scatter(X[:,0], X[:,1], c=kmeans.labels_)
ax.set_title("K-Means Clustering Results with K=2");
```



Next we try and cluster with K=3.

```
[ ]: kmeans = KMeans(n_clusters=3, random_state=3,  
n_init=20).fit(X)  
fig, ax = plt.subplots(figsize=(8,8))  
ax.scatter(X[:,0], X[:,1], c=kmeans.labels_)  
ax.set_title("K-Means Clustering Results with K=3");
```



```
[ ]: kmeans1 = KMeans(n_clusters=3, random_state=3,  
n_init=1).fit(X)  
kmeans20 = KMeans(n_clusters=3, random_state=3,  
n_init=20).fit(X);  
kmeans1.inertia_, kmeans20.inertia_
```

```
[ ]: (76.85131986999252, 75.06261242745384)
```

The inertia (which can be calculated with `.inertia_` gives is the total within-cluster sum of squares. We seek to minimize this value with K-means clustering.

Some steps to take when we perform K-means are - Use random multiple initial cluster assignments
- Use a random seed in the `random_state` argument for `Kmeans()` This is important so that you

can replicate the first step of the algorithm