

lab3

October 27, 2023

0.1 Lab 3: Linear Regression

Jack Krebsbach Math 313

```
[ ]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
import statsmodels.api as sm
```

Here we import some submodules from the `statsmodels` packages. We rename the named exports to keep the *namespace* clean.

```
[ ]: from statsmodels.stats.outliers_influence import variance_inflation_factor as VIF
from statsmodels.stats.anova import anova_lm
```

We also use some functions from ISLP

```
[ ]: from ISLP import load_data
from ISLP.models import (ModelSpec as MS, summarize, poly)
```

0.1.1 Objects & Namespaces

We can list the objects in the current namespace using `dir()` (the following is truncated).

```
[ ]: dir()[:5]
```

```
[ ]: ['A', 'Boston', 'Carseats', 'In', 'MS']
```

Each python object has its own namespace. We can also access this namespace of specific objects in the following way. Notice that `numpy` arrays has a `sum` property!

```
[ ]: A = np.arange(0,10,2)
dir(A)[:5]
```

```
[ ]: ['T', '__abs__', '__add__', '__and__', '__array__']
```

```
[ ]: A.sum()
```

```
[ ]: 20
```

0.1.2 Simple Linear Regression

First we will create design/model matrices using `ModelSpec()` transform from `ISLP.models`.

The task will be to build a model using 13 predictors to predict `medv`.

```
[ ]: Boston = load_data("Boston")
      Boston.columns

[ ]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
          'ptratio', 'lstat', 'medv'],
          dtype='object')
```

We can learn more about the data by typing `Boston?`. Our first model will just use a single predictor to predict `medv`.

```
[ ]: Boston?

[ ]: # The model matrix is constructed by hand
      X = pd.DataFrame({'intercept': np.ones(Boston.shape[0]), 'lstat':
          ↪Boston['lstat']})

      X[:4]
```

```
[ ]:      intercept  lstat
      0          1.0   4.98
      1          1.0   9.14
      2          1.0   4.03
      3          1.0   2.94
```

```
[ ]: # Here we fit the model
      y = Boston['medv']
      model = sm.OLS(y, X)
      results = model.fit()
```

To get more information about the fitted model we can use the `summarize` function. Which gives useful statistics like standard errors, t-statistics, and p-values.

```
[ ]: summarize(results)

[ ]:      coef  std err      t  P>|t|
      intercept  34.5538    0.563  61.415    0.0
      lstat      -0.9500    0.039 -24.528    0.0
```

Using Transformations: Fit and Transform In practice models usually contain more than one predictor. In addition, transformations of the variables and interactions terms can be added. We can rely on tools from `sklearn` to create these transforms.

```
[ ]: design = MS(['lstat'])
design = design.fit(Boston)
X = design.transform(Boston)
X[:4]
```

```
[ ]:      intercept  lstat
0         1.0     4.98
1         1.0     9.14
2         1.0     4.03
3         1.0     2.94
```

In our previous data set `fit()` does not do much, it just checks if the `lstat` variable exists in the data set.

While this processed was executed in two lines of code, the design object is changed after calling `fit()`.

```
[ ]: results.summary()
results.params
```

```
[ ]: intercept    34.553841
lstat           -0.950049
dtype: float64
```

We can also use `get_prediction` to obtain new labels with confidence intervals from data not in the training set.

```
[ ]: new_df = pd.DataFrame({'lstat':[5, 10, 15]})
newX = design.transform(new_df)
newX
```

```
[ ]:      intercept  lstat
0         1.0       5
1         1.0      10
2         1.0      15
```

```
[ ]: # Here we obtain predictions for the new data.
new_predictions = results.get_prediction(newX)
new_predictions.predicted_mean
```

```
[ ]: array([29.80359411, 25.05334734, 20.30310057])
```

```
[ ]: # Extract confidence intervals
new_predictions.conf_int(alpha=0.05)
```

```
[ ]: array([[29.00741194, 30.59977628],
          [24.47413202, 25.63256267],
          [19.73158815, 20.87461299]])
```

```
[ ]: # Extract prediction intervals
new_predictions.conf_int(obs=True, alpha=0.05)
```

```
[ ]: array([[17.56567478, 42.04151344],
          [12.82762635, 37.27906833],
          [ 8.0777421 , 32.52845905]])
```

As expected the prediction intervals are centered around the same values as the confidence intervals, but are wider. This is because the confidence interval relays information on the *average* of our data, while the prediction interval gives an information on the confidence for a new, single city, which incorporates the irreducible error in addition to the reducible error.


Defining functions Here we define some helper functions to visualize the fits and residuals.

```
[ ]: def abline(ax, b, m):
      "Add a line with slope m and intercept b to ax"
      xlim = ax.get_xlim()
      ylim = [m * xlim[0] + b, m * xlim[1] + b]
      ax.plot(xlim, ylim)
```

By adding in `*args` and `**kwargs` we can any number of non-named arguments and any number of named arguments.

```
[ ]: def abline(ax, b, m, *args, **kwargs):
      "Add a line with slope m and intercept b to ax"
      xlim = ax.get_xlim()
      ylim = [m * xlim[0] + b, m * xlim[1] + b]
      ax.plot(xlim, ylim, *args, **kwargs)
```


```
[ ]: ax = Boston.plot.scatter('lstat', 'medv')
      abline(ax,
              results.params[0],
              results.params[1],
              'r--', linewidth=3)
```



The figure above shows that there is some evidence of a non-linear relationship between `lstat` and `medv`.


Here we analyze plots that visualize information of the fitted values and residuals. The residual plot also shows evidence of a non-linear relationship. We can also compute leverage statistics by using the `hat_matrix_diag` attribute from the object returned by `get_influence()`.

```
[ ]: ax = subplots(figsize=(8,8))[1]
      ax.scatter(results.fittedvalues , results.resid)
      ax.set_xlabel('Fitted value')
      ax.set_ylabel('Residual')
      ax.axhline(0, c='k', ls='--');
```



```
[ ]: infl = results.get_influence()
ax = subplots(figsize=(8,8))[1]
ax.scatter(np.arange(X.shape[0]), infl.hat_matrix_diag)
ax.set_xlabel('Index')
ax.set_ylabel('Leverage')
```

```
[ ]: Text(0, 0.5, 'Leverage')
```



```
[ ]: # This identifies the index of the largest element in the array which in turn  

     ↪ identifies which observation has the largest leverage statistic.  

     np.argmax(infl.hat_matrix_diag)
```

```
[ ]: 374
```

0.1.3 Multiple Linear Regression

To use multiple linear regression using least squares we can use `ModelSpec()` transform to construct the required model matrix. In the following we add `age` to the predictors.

```
[ ]: X = MS(['lstat', 'age']).fit_transform(Boston)  

     model1 = sm.OLS(y, X)  

     results1 = model1.fit()
```

```
summarize(results1)
```

```
[ ]:      coef  std err      t  P>|t|
intercept  33.2228    0.731  45.458  0.000
lstat      -1.0321    0.048 -21.416  0.000
age         0.0345    0.012   2.826  0.005
```

We can easily create a list of all the predictors by dropping a single one instead of typing them all out.

```
[ ]: terms = Boston.columns.drop('medv')
terms
```

```
[ ]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
          'ptratio', 'lstat'],
          dtype='object')
```

```
[ ]: # Fit the multiple linear regression model.
X = MS(terms).fit_transform(Boston)
model = sm.OLS(y, X)
results = model.fit()
summarize(results)
```

```
[ ]:      coef  std err      t  P>|t|
intercept  41.6173    4.936   8.431  0.000
crim       -0.1214    0.033  -3.678  0.000
zn         0.0470    0.014   3.384  0.001
indus      0.0135    0.062   0.217  0.829
chas       2.8400    0.870   3.264  0.001
nox       -18.7580    3.851  -4.870  0.000
rm         3.6581    0.420   8.705  0.000
age        0.0036    0.013   0.271  0.787
dis       -1.4908    0.202  -7.394  0.000
rad        0.2894    0.067   4.325  0.000
tax       -0.0127    0.004  -3.337  0.001
ptratio   -0.9375    0.132  -7.091  0.000
lstat     -0.5520    0.051 -10.897  0.000
```

```
[ ]: # Using all the variables as predictors except age
minus_age = Boston.columns.drop(['medv', 'age'])
Xma = MS(minus_age).fit_transform(Boston)
model1 = sm.OLS(y, Xma)
summarize(model1.fit())
```

```
[ ]:      coef  std err      t  P>|t|
intercept  41.5251    4.920   8.441  0.000
crim       -0.1214    0.033  -3.683  0.000
zn         0.0465    0.014   3.379  0.001
```


indus	0.0135	0.062	0.217	0.829
chas	2.8528	0.868	3.287	0.001
nox	-18.4851	3.714	-4.978	0.000
rm	3.6811	0.411	8.951	0.000
dis	-1.5068	0.193	-7.825	0.000
rad	0.2879	0.067	4.322	0.000
tax	-0.0127	0.004	-3.333	0.001
ptratio	-0.9346	0.132	-7.099	0.000
lstat	-0.5474	0.048	-11.483	0.000

Multivariable goodness of fit. The individual components of results can be accessed by name.

To access this information we can use list comprehension which is a simple and powerful way to form a list in python.

```
[ ]: # Forming these Python objects within the list is called list comprehension.
vals = [VIF(X, i)
        for i in range(1, X.shape[1])]
vif = pd.DataFrame({'vif':vals},
                    index=X.columns[1:])
vif
```

```
[ ]:
      crim      1.767486
      zn       2.298459
      indus    3.987181
      chas     1.071168
      nox      4.369093
      rm       1.912532
      age      3.088232
      dis      3.954037
      rad      7.445301
      tax      9.002158
      ptratio  1.797060
      lstat    2.870777
```

```
[ ]: # We can use list comprehension to perform repetitive operations
vals = []
for i in range(1, X.values.shape[1]):
    vals.append(VIF(X.values, i))
```

Interaction terms To add an interaction terms we can include a tuple in the model matrix.

```
[ ]: X = MS(['lstat', 'age', ('lstat', 'age')]).fit_transform(Boston)
model2 = sm.OLS(y, X)
summarize(model2.fit())
```

```
[ ]:      coef  std err      t  P>|t|
intercept  36.0885    1.470  24.553  0.000
lstat      -1.3921    0.167  -8.313  0.000
age        -0.0007    0.020  -0.036  0.971
lstat:age   0.0042    0.002   2.244  0.025
```

Non-linear Transformation of the Predictors We can also include non-linear transformations of the predictors in addition to just the interaction terms and features themselves.

```
[ ]: X = MS([poly('lstat', degree=2), 'age']).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)
```

```
[ ]:      coef  std err      t  P>|t|
intercept      17.7151    0.781  22.681    0.0
poly(lstat, degree=2)[0] -179.2279    6.733 -26.620    0.0
poly(lstat, degree=2)[1]   72.9908    5.482  13.315    0.0
age              0.0703    0.011   6.471    0.0
```

We see that the p-value associated with the quadratic term is near zero. This means that the model was improved by adding that term.

```
[ ]: anova_lm(results1, results3)
```

```
[ ]:      df_resid      ssr df_diff      ss_diff      F      Pr(>F)
0      503.0  19168.128609    0.0      NaN      NaN      NaN
1      502.0  14165.613251    1.0  5002.515357  177.278785  7.468491e-35
```


To get a better understanding of how the quadratic fit is superior to the linear fit we can use the `anova_lm` function, which performs a hypothesis test against the two models.

The F statistics is large and its corresponding p-value is near zero, suggesting ample evidence that the bigger model is better.

In the figure below we see that the residuals do not have a discernible structure or pattern. Therefore, we can be confident that our model can appropriately describe the underlying pattern of the data.

```
[ ]: ax = subplots(figsize=(8,8))[1]
ax.scatter(results3.fittedvalues , results3.resid)
ax.set_xlabel('Fitted value')
ax.set_ylabel('Residual')
ax.axhline(0, c='k', ls='--')
```

```
[ ]: <matplotlib.lines.Line2D at 0x169535c10>
```



Qualitative Predictors Instead of the previous data set we will work with the `Carseats` data which contains qualitative predictors in addition to quantitative predictors. We will work to predict `sales` in 400 locations.

These data are realized as one-hot-encoded variables.

```
[ ]: Carseats = load_data('Carseats')
      Carseats.columns

[ ]: Index(['Sales', 'CompPrice', 'Income', 'Advertising', 'Population', 'Price',
           'ShelveLoc', 'Age', 'Education', 'Urban', 'US'],
          dtype='object')
```

To avoid collinearity we drop the first column.

```
[ ]: allvars = list(Carseats.columns.drop('Sales'))
y = Carseats['Sales']
final = allvars + [('Income', 'Advertising'), ('Price', 'Age')]
X = MS(final).fit_transform(Carseats)
model = sm.OLS(y, X)
summarize(model.fit())
```

```
[ ]:
      coef  std err      t  P>|t|
intercept      6.5756    1.009    6.519  0.000
CompPrice      0.0929    0.004   22.567  0.000
Income         0.0109    0.003    4.183  0.000
Advertising     0.0702    0.023    3.107  0.002
Population     0.0002    0.000    0.433  0.665
Price        -0.1008    0.007  -13.549  0.000
ShelveLoc[Good]  4.8487    0.153   31.724  0.000
ShelveLoc[Medium] 1.9533    0.126   15.531  0.000
Age          -0.0579    0.016   -3.633  0.000
Education     -0.0209    0.020   -1.063  0.288
Urban[Yes]      0.1402    0.112    1.247  0.213
US[Yes]        -0.1576    0.149   -1.058  0.291
Income:Advertising 0.0008    0.000    2.698  0.007
Price:Age       0.0001    0.000    0.801  0.424
```

We have added an interaction term between price and age. To encode the `ShelveLoc[Good]` dummy variable we insert a 1 to indicate a positive observation and a 0 otherwise.

We see that `ShelveLoc[Good]` has a positive value, indicating that good shelving location is associated with high sales. It has a higher coefficient than `ShelveLoc[Medium]` which means it leads to higher sales as well.