

# lab4

November 28, 2023

## 0.1 Lab 4: Classification

Jack Krebsbach Math 313

### Imports

```
[ ]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import (ModelSpec as MS,
summarize)
```

### New imports needed

```
[ ]: from ISLP import confusion_table
from ISLP.models import contrast
from sklearn.discriminant_analysis import \
(LinearDiscriminantAnalysis as LDA, QuadraticDiscriminantAnalysis as QDA)
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
[ ]: # Load the data
Smarket = load_data('Smarket')
Smarket
```

```
[ ]:
      Year  Lag1  Lag2  Lag3  Lag4  Lag5  Volume  Today  Direction
0      2001  0.381 -0.192 -2.624 -1.055  5.010  1.19130  0.959         Up
1      2001  0.959  0.381 -0.192 -2.624 -1.055  1.29650  1.032         Up
2      2001  1.032  0.959  0.381 -0.192 -2.624  1.41120 -0.623        Down
3      2001 -0.623  1.032  0.959  0.381 -0.192  1.27600  0.614         Up
4      2001  0.614 -0.623  1.032  0.959  0.381  1.20570  0.213         Up
...
1245  2005  0.422  0.252 -0.024 -0.584 -0.285  1.88850  0.043         Up
1246  2005  0.043  0.422  0.252 -0.024 -0.584  1.28581 -0.955        Down
```

1247	2005	-0.955	0.043	0.422	0.252	-0.024	1.54047	0.130	Up
1248	2005	0.130	-0.955	0.043	0.422	0.252	1.42236	-0.298	Down
1249	2005	-0.298	0.130	-0.955	0.043	0.422	1.38254	-0.489	Down

[1250 rows x 9 columns]

```
[ ]: # Columns of data set
Smarket.columns
```

```
[ ]: Index(['Year', 'Lag1', 'Lag2', 'Lag3', 'Lag4', 'Lag5', 'Volume', 'Today',
          'Direction'],
          dtype='object')
```

```
[ ]: Smarket.corr()
```


```
/var/folders/gf/bt25hkv172n_bttx0h72_6340000gn/T/ipykernel_12419/1422385858.py:1
: FutureWarning: The default value of numeric_only in DataFrame.corr is
deprecated. In a future version, it will default to False. Select only valid
columns or specify the value of numeric_only to silence this warning.
Smarket.corr()
```

```
[ ]:
      Year      Lag1      Lag2      Lag3      Lag4      Lag5      Volume \
Year      1.000000  0.029700  0.030596  0.033195  0.035689  0.029788  0.539006
Lag1      0.029700  1.000000 -0.026294 -0.010803 -0.002986 -0.005675  0.040910
Lag2      0.030596 -0.026294  1.000000 -0.025897 -0.010854 -0.003558 -0.043383
Lag3      0.033195 -0.010803 -0.025897  1.000000 -0.024051 -0.018808 -0.041824
Lag4      0.035689 -0.002986 -0.010854 -0.024051  1.000000 -0.027084 -0.048414
Lag5      0.029788 -0.005675 -0.003558 -0.018808 -0.027084  1.000000 -0.022002
Volume     0.539006  0.040910 -0.043383 -0.041824 -0.048414 -0.022002  1.000000
Today      0.030095 -0.026155 -0.010250 -0.002448 -0.006900 -0.034860  0.014592

      Today
Year      0.030095
Lag1     -0.026155
Lag2     -0.010250
Lag3     -0.002448
Lag4     -0.006900
Lag5     -0.034860
Volume    0.014592
Today     1.000000
```

```
[ ]: # Volume is increasing over time.
Smarket.plot(y='Volume')
```

```
[ ]: <Axes: >
```



**Logistic Regression** We will fit a logistic regression models to predict Direction using Lag1 through Lag5.

```
[ ]: # To run LR we use family = sm.families.Binomial()
allvars = Smarket.columns.drop(['Today', 'Direction', 'Year'])
design = MS(allvars)
X = design.fit_transform(Smarket)
y = Smarket.Direction == 'Up'
glm = sm.GLM(y, X, family=sm.families.Binomial())
results = glm.fit()
summarize(results)
```

```
[ ]:
```

	coef	std err	z	P> z
intercept	-0.1260	0.241	-0.523	0.601
Lag1	-0.0731	0.050	-1.457	0.145
Lag2	-0.0423	0.050	-0.845	0.398
Lag3	0.0111	0.050	0.222	0.824
Lag4	0.0094	0.050	0.187	0.851
Lag5	0.0103	0.050	0.208	0.835
Volume	0.1354	0.158	0.855	0.392

The smallest p-value here is associated with Lag1, but it is still not enough to provide clear evidence

of an association between `Lag1` and `Direction`. Because it is a negative coefficient that means a positive increase in market in the previous day suggests it is less likely to go up today.

```
[ ]: # Grab the coefficients
      results.params
```

```
[ ]: intercept    -0.126000
      Lag1         -0.073074
      Lag2         -0.042301
      Lag3          0.011085
      Lag4          0.009359
      Lag5          0.010313
      Volume       0.135441
      dtype: float64
```

```
[ ]: # Grab the p-values
      results.pvalues
```

```
[ ]: intercept    0.600700
      Lag1         0.145232
      Lag2         0.398352
      Lag3         0.824334
      Lag4         0.851445
      Lag5         0.834998
      Volume       0.392404
      dtype: float64
```

```
[ ]: probs = results.predict()
      probs[:10]
```

```
[ ]: array([0.50708413, 0.48146788, 0.48113883, 0.51522236, 0.51078116,
            0.50695646, 0.49265087, 0.50922916, 0.51761353, 0.48883778])
```

To predict a binary response – up or down – we must first convert these probabilities to class labels based on if the probability is greater or less than 0.5.

```
[ ]: # Create array of length 1250 with each element Down
      labels = np.array(['Down']*1250)
      labels[probs>0.5] = "Up"
```

```
[ ]: confusion_table(labels, Smarket.Direction)
```

```
[ ]: Truth      Down  Up
      Predicted
      Down      145  141
      Up        457  507
```

```
[ ]: (507+145)/1250, np.mean(labels == Smarket.Direction)
```

```
[ ]: (0.5216, 0.5216)
```

The training error rate here is  $100 - 52.2 = 47.8$ , which is overly optimistic. To get a more clear idea of the true error rate we create a training and testing set.

```
[ ]: train = (Smarket.Year < 2005)
Smarket_train = Smarket.loc[train]
Smarket_test = Smarket.loc[~train]
Smarket_test.shape
```

```
[ ]: (252, 9)
```

```
[ ]: X_train, X_test = X.loc[train], X.loc[~train]
y_train, y_test = y.loc[train], y.loc[~train]
glm_train = sm.GLM(y_train, X_train, family=sm.families.Binomial())
results = glm_train.fit()
probs = results.predict(exog=X_test)
```

```
[ ]: D = Smarket.Direction
L_train, L_test = D.loc[train], D.loc[~train]
```

```
[ ]: # Create array for all observations
labels = np.array(['Down']*252)
# Classify with decision rule
labels[probs>0.5] = 'Up'
# Create confusion matrix
confusion_table(labels, L_test)
```

```
[ ]: Truth      Down  Up
Predicted
Down          77  97
Up            34  44
```

```
[ ]: np.mean(labels == L_test), np.mean(labels != L_test)
```

```
[ ]: (0.4801587301587302, 0.5198412698412699)
```

- The diagonals of the decision matrix are the correct values.
- The test accuracy is ~48%
- The test error rate is ~52%

```
[ ]: # We build a model with just Lag1 and Lag2
model = MS(['Lag1', 'Lag2']).fit(Smarket)
X = model.transform(Smarket)
X_train, X_test = X.loc[train], X.loc[~train]
glm_train = sm.GLM(y_train, X_train, family=sm.families.Binomial())
results = glm_train.fit()
probs = results.predict(exog=X_test)
```

```
labels = np.array(['Down']*252)
labels[probs>0.5] = 'Up'
confusion_table(labels, L_test)
```

```
[ ]: Truth      Down   Up
      Predicted
      Down      35    35
      Up       76   106
```

```
[ ]: (35+106) / 252 , 106/(106+76)
```

```
[ ]: (0.5595238095238095, 0.5824175824175825)
```

```
[ ]: newdata = pd.DataFrame({'Lag1':[1.2, 1.5], 'Lag2':[1.1, -0.8]});
```

```
[ ]: newX = model.transform(newdata)
      results.predict(newX)
```

```
[ ]: 0    0.479146
      1    0.496094
      dtype: float64
```

We see that when in this new model the overall accuracy goes up. This logistic regression model has a 58% accuracy rate predicting increases in the market.

**Linear Discriminant Analysis** We perform Linear Discriminant Analysis on the Smarket data.

```
[ ]: # Create instance of LDA
      lda = LDA(store_covariance=True)
```

```
[ ]: # Split the train ant test set
      X_train, X_test = [M.drop(columns=['intercept']) for M in [X_train, X_test]]
      lda.fit(X_train, L_train)
```

```
[ ]: LinearDiscriminantAnalysis(store_covariance=True)
```

We can extract the means of the two classes with the means attribute on the lda object.

```
[ ]: lda.means_
```

```
[ ]: array([[ 0.04279022,  0.03389409],
            [-0.03954635, -0.03132544]])
```

To insure we know which class is corresponding to which label with lda.classes\_.

```
[ ]: lda.classes_
```

```
[ ]: array(['Down', 'Up'], dtype='<U4')
```

Extracting the priors yeilds

```
[ ]: lda.priors_  
[ ]: array([0.49198397, 0.50801603])
```

```
[ ]: # Make predictions on the test set  
lda_pred = lda.predict(X_test)
```

```
[ ]: # The linear discriminant vectors  
lda.scalings_
```

```
[ ]: array([[ -0.64201904],  
          [ -0.51352928]])
```

```
[ ]: # Make predictions  
lda_pred = lda.predict(X_test)
```

```
[ ]: # Create confusing matrix  
confusion_table(lda_pred, L_test)
```

```
[ ]: Truth      Down   Up  
Predicted  
Down          35    35  
Up            76   106
```

```
[ ]: lda_prob = lda.predict_proba(X_test)  
np.all(  
np.where(lda_prob[:,1] >= 0.5, 'Up', 'Down') == lda_pred )
```

```
[ ]: True
```

We can also use the posterior probabilities with a 50% threshold to re-create the predictions from the fitted lda instance.

```
[ ]: np.all(  
[lda.classes_[i] for i in np.argmax(lda_prob, 1)] ==  
lda_pred )
```

```
[ ]: True
```

```
[ ]: np.sum(lda_prob[:,0] > 0.9)
```

```
[ ]: 0
```

Interestingly, there are no days that meet the threshold of 90%!

**Quadratic Discriminant Analysis** Fit a QDA model on the Smarket data

```
[ ]: # Instantiate the model
qda = QDA(store_covariance=True)
qda.fit(X_train, L_train)
```

```
[ ]: QuadraticDiscriminantAnalysis(store_covariance=True)
```

```
[ ]: # Get the means and priors
qda.means_, qda.priors_
```

```
[ ]: (array([[ 0.04279022,  0.03389409],
            [-0.03954635, -0.03132544]]),
      array([0.49198397, 0.50801603]))
```

The QDA() function will compute each a covariance matrix for each class.

```
[ ]: # Extract the covariance matrix for the first class.
qda.covariance_[0]
```

```
[ ]: array([[ 1.50662277, -0.03924806],
            [-0.03924806,  1.53559498]])
```

Just like before we can make predictions on the test set using the trained classifier.

```
[ ]: qda_pred = qda.predict(X_test)
confusion_table(qda_pred, L_test)
```

```
[ ]: Truth      Down   Up
Predicted
Down         30    20
Up           81   121
```

```
[ ]: np.mean(qda_pred == L_test)
```

```
[ ]: 0.5992063492063492
```

An accuracy this high for stock market data suggests that QDA has the possibility of capturing the true relationship more than the linear forms.

## Naive Bayes

```
[ ]: # Instantiate object
NB = GaussianNB()
NB.fit(X_train, L_train)
```

```
[ ]: GaussianNB()
```

```
[ ]: # Get the classes
NB.classes_
```



```
[ ]: array(['Down', 'Up'], dtype='<U4')
```

```
[ ]: # Extract the prior probabilities  
NB.class_prior_
```

```
[ ]: array([0.49198397, 0.50801603])
```

```
[ ]: # These are the means for each fitted model by each class and feature  
NB.theta_
```

```
[ ]: array([[ 0.04279022,  0.03389409],  
          [-0.03954635, -0.03132544]])
```

```
[ ]: # Similarly here are the variances. 2 classes * 2 features= 4 variances  
NB.var_
```

```
[ ]: array([[1.50355429, 1.53246749],  
          [1.51401364, 1.48732877]])
```

```
[ ]: mean= X_train[L_train == 'Down'].mean()  
      variance = X_train[L_train == 'Down'].var(ddof=0)  
      print(f'Mean: {mean}')  
      print(f'Variance: {variance}')
```

```
Mean: Lag1    0.042790  
Lag2    0.033894  
dtype: float64  
Variance: Lag1    1.503554  
Lag2    1.532467  
dtype: float64
```

```
[ ]: # Extract confusion matrix  
nb_labels = NB.predict(X_test)  
confusion_table(nb_labels, L_test)
```

```
[ ]: Truth      Down   Up  
      Predicted  
Down      29   20  
Up       82  121
```

```
[ ]: # Predict on new data points  
NB.predict_proba(X_test)[:5]
```

```
[ ]: array([[0.4873288 , 0.5126712 ],  
          [0.47623584, 0.52376416],  
          [0.46529531, 0.53470469],  
          [0.47484469, 0.52515531],
```

```
[0.49020587, 0.50979413]])
```

Using `NB.predict_proba` we can estimate the probability that each observation belongs to a particular class.

**K-Nearest Neighbors** We can also use K-Nearest Neighbors to create a classifier

```
[ ]: knn1 = KNeighborsClassifier(n_neighbors=1)
      knn1.fit(X_train, L_train)
      knn1_pred = knn1.predict(X_test)
      confusion_table(knn1_pred, L_test)
```

```
[ ]: Truth      Down  Up
      Predicted
      Down      43   58
      Up        68   83
```

```
[ ]: # Not a very good fit with 50 percent accuracy
      (83+43)/252, np.mean(knn1_pred == L_test)
```

```
[ ]: (0.5, 0.5)
```

Since we did not get a good accuracy with 2 nearest neighbors we can try with 3.

```
[ ]: knn3 = KNeighborsClassifier(n_neighbors=3)
      knn3_pred = knn3.fit(X_train, L_train).predict(X_test)
      np.mean(knn3_pred == L_test)
```

```
[ ]: 0.5317460317460317
```

We only did a little bit better, now with 53% accuracy. Thus, KNN does not do well on the `Smarket` data set. We can see its utility using the `Caravan` data set.

```
[ ]: Caravan = load_data('Caravan')
      Purchase = Caravan.Purchase
      Purchase.value_counts()
```

```
[ ]: No      5474
      Yes     348
      Name: Purchase, dtype: int64
```

```
[ ]: # Proportion of individuals that purchase a caravan insurance policy.
      348 / 5822
```

```
[ ]: 0.05977327378907592
```

```
[ ]: # Create the feature data frame
      feature_df = Caravan.drop(columns=['Purchase'])
```

```
[ ]: # Create scaler
scaler = StandardScaler(with_mean=True, with_std=True,
copy=True)
```

When argument with\_mean is True the means are subtracted off.

```
[ ]: # Scale the data
scaler.fit(feature_df)
X_std = scaler.transform(feature_df)
```

```
[ ]: feature_std = pd.DataFrame( X_std , columns=feature_df.columns);
feature_std.std()
```

```
[ ]: MOSTYPE      1.000086
MAANTHUI      1.000086
MGEMOMV       1.000086
MGEMLEEF      1.000086
MOSHOOFD      1.000086
...
AZEILPL       1.000086
APLEZIER      1.000086
AFIETS        1.000086
AINBOED       1.000086
ABYSTAND      1.000086
Length: 85, dtype: float64
```

```
[ ]: (X_train, X_test , y_train , y_test) = train_test_split(feature_std, Purchase ,
↳test_size=1000, random_state=0)
```

```
[ ]: # Instantiate classifier
knn1 = KNeighborsClassifier(n_neighbors=1)
# Fit the model
knn1_pred = knn1.fit(X_train, y_train).predict(X_test)
# Get the fitted values.
np.mean(y_test != knn1_pred), np.mean(y_test != "No")
```

```
[ ]: (0.111, 0.067)
```

```
[ ]: # Confusion matrix
confusion_table(knn1_pred, y_test)
```

```
[ ]: Truth      No  Yes
Predicted
No          880  58
Yes         53   9
```

```
[ ]: 9/(53+9)
```

```
[ ]: 0.14516129032258066
```

The KNN model with K=1 does a good job at predicting customers that buy insurance. This is double what you would get if you were just guessing.

### Tuning Parameters

```
[ ]: for K in range(1,6):
    knn = KNeighborsClassifier(n_neighbors=K)
    knn_pred = knn.fit(X_train, y_train).predict(X_test)
    C = confusion_table(knn_pred, y_test)
    templ = ('K={0:d}: # predicted to rent: {1:>2}, ' +
            ' # who did rent {2:d}, accuracy {3:.1%}')
```

```
    pred = C.loc['Yes'].sum()
    did_rent = C.loc['Yes', 'Yes']
    print(templ.format(K, pred , did_rent ,did_rent / pred))
```

```
K=1: # predicted to rent: 62, # who did rent 9, accuracy 14.5%
K=2: # predicted to rent:  6, # who did rent 1, accuracy 16.7%
K=3: # predicted to rent: 20, # who did rent 3, accuracy 15.0%
K=4: # predicted to rent:  4, # who did rent 0, accuracy 0.0%
K=5: # predicted to rent:  7, # who did rent 1, accuracy 14.3%
```

Here we see that at K=4 the accuracy and predictions is very different than the rest.

### Comparison to Logistic Regression

```
[ ]: logit = LogisticRegression(C=1e10, solver='liblinear')
    logit.fit(X_train, y_train)
    logit_pred = logit.predict_proba(X_test)
    logit_labels = np.where(logit_pred[:,1] > 5, 'Yes', 'No')
    confusion_table(logit_labels, y_test)
```

```
[ ]: Truth      No  Yes
    Predicted
    No          933  67
    Yes          0   0
```

```
[ ]: logit_labels = np.where(logit_pred[:,1]>0.25, 'Yes', 'No')
    confusion_table(logit_labels, y_test)
```

```
[ ]: Truth      No  Yes
    Predicted
    No          913  58
    Yes          20   9
```

```
[ ]: 9/(20+9)
```

```
[ ]: 0.3103448275862069
```

## Linear and Poisson Regression on Bikeshare Data

```
[ ]: # Load the data
      Bike = load_data('Bikeshare')
```

First we fit a linear regression model to the data.

```
[ ]: X = MS(['mnth', 'hr',
            'workingday', 'temp',
            'weathersit']).fit_transform(Bike)
      Y = Bike['bikers']
      M_lm = sm.OLS(Y, X).fit()
      summarize(M_lm)
```

```
[ ]:
```

	coef	std err	t	P> t
intercept	-68.6317	5.307	-12.932	0.000
mnth[Feb]	6.8452	4.287	1.597	0.110
mnth[March]	16.5514	4.301	3.848	0.000
mnth[April]	41.4249	4.972	8.331	0.000
mnth[May]	72.5571	5.641	12.862	0.000
mnth[June]	67.8187	6.544	10.364	0.000
mnth[July]	45.3245	7.081	6.401	0.000
mnth[Aug]	53.2430	6.640	8.019	0.000
mnth[Sept]	66.6783	5.925	11.254	0.000
mnth[Oct]	75.8343	4.950	15.319	0.000
mnth[Nov]	60.3100	4.610	13.083	0.000
mnth[Dec]	46.4577	4.271	10.878	0.000
hr[1]	-14.5793	5.699	-2.558	0.011
hr[2]	-21.5791	5.733	-3.764	0.000
hr[3]	-31.1408	5.778	-5.389	0.000
hr[4]	-36.9075	5.802	-6.361	0.000
hr[5]	-24.1355	5.737	-4.207	0.000
hr[6]	20.5997	5.704	3.612	0.000
hr[7]	120.0931	5.693	21.095	0.000
hr[8]	223.6619	5.690	39.310	0.000
hr[9]	120.5819	5.693	21.182	0.000
hr[10]	83.8013	5.705	14.689	0.000
hr[11]	105.4234	5.722	18.424	0.000
hr[12]	137.2837	5.740	23.916	0.000
hr[13]	136.0359	5.760	23.617	0.000
hr[14]	126.6361	5.776	21.923	0.000
hr[15]	132.0865	5.780	22.852	0.000
hr[16]	178.5206	5.772	30.927	0.000
hr[17]	296.2670	5.749	51.537	0.000
hr[18]	269.4409	5.736	46.976	0.000
hr[19]	186.2558	5.714	32.596	0.000
hr[20]	125.5492	5.704	22.012	0.000
hr[21]	87.5537	5.693	15.378	0.000

hr[22]	59.1226	5.689	10.392	0.000
hr[23]	26.8376	5.688	4.719	0.000
workingday	1.2696	1.784	0.711	0.477
temp	157.2094	10.261	15.321	0.000
weathersit[cloudy/misty]	-12.8903	1.964	-6.562	0.000
weathersit[heavy rain/snow]	-109.7446	76.667	-1.431	0.152
weathersit[light rain/snow]	-66.4944	2.965	-22.425	0.000

We see that there are 24 levels and 40 observations.

```
[ ]: # We change the encoding of the variables hr and mnth
```

```
hr_encode = contrast('hr', 'sum')
mnth_encode = contrast('mnth', 'sum')
```

```
[ ]: X2 = MS([mnth_encode, hr_encode, 'workingday', 'temp', 'weathersit']).
```

```
    fit_transform(Bike)
M2_lm = sm.OLS(Y, X2).fit()
S2 = summarize(M2_lm)
S2
```

```
[ ]:
```

	coef	std err	t	P> t
intercept	73.5974	5.132	14.340	0.000
mnth[Jan]	-46.0871	4.085	-11.281	0.000
mnth[Feb]	-39.2419	3.539	-11.088	0.000
mnth[March]	-29.5357	3.155	-9.361	0.000
mnth[April]	-4.6622	2.741	-1.701	0.089
mnth[May]	26.4700	2.851	9.285	0.000
mnth[June]	21.7317	3.465	6.272	0.000
mnth[July]	-0.7626	3.908	-0.195	0.845
mnth[Aug]	7.1560	3.535	2.024	0.043
mnth[Sept]	20.5912	3.046	6.761	0.000
mnth[Oct]	29.7472	2.700	11.019	0.000
mnth[Nov]	14.2229	2.860	4.972	0.000
hr[0]	-96.1420	3.955	-24.307	0.000
hr[1]	-110.7213	3.966	-27.916	0.000
hr[2]	-117.7212	4.016	-29.310	0.000
hr[3]	-127.2828	4.081	-31.191	0.000
hr[4]	-133.0495	4.117	-32.319	0.000
hr[5]	-120.2775	4.037	-29.794	0.000
hr[6]	-75.5424	3.992	-18.925	0.000
hr[7]	23.9511	3.969	6.035	0.000
hr[8]	127.5199	3.950	32.284	0.000
hr[9]	24.4399	3.936	6.209	0.000
hr[10]	-12.3407	3.936	-3.135	0.002
hr[11]	9.2814	3.945	2.353	0.019
hr[12]	41.1417	3.957	10.397	0.000
hr[13]	39.8939	3.975	10.036	0.000

hr[14]	30.4940	3.991	7.641	0.000
hr[15]	35.9445	3.995	8.998	0.000
hr[16]	82.3786	3.988	20.655	0.000
hr[17]	200.1249	3.964	50.488	0.000
hr[18]	173.2989	3.956	43.806	0.000
hr[19]	90.1138	3.940	22.872	0.000
hr[20]	29.4071	3.936	7.471	0.000
hr[21]	-8.5883	3.933	-2.184	0.029
hr[22]	-37.0194	3.934	-9.409	0.000
workingday	1.2696	1.784	0.711	0.477
temp	157.2094	10.261	15.321	0.000
weathersit[cloudy/misty]	-12.8903	1.964	-6.562	0.000
weathersit[heavy rain/snow]	-109.7446	76.667	-1.431	0.152
weathersit[light rain/snow]	-66.4944	2.965	-22.425	0.000

Overall, we see that the choice of coding does not really matter, as long as we interpret the model output correctly.

```
[ ]: # The sum of the squared differences is zero
np.sum((M_lm.fittedvalues - M2_lm.fittedvalues)**2)
```

```
[ ]: 5.006155854534554e-20
```

```
[ ]: np.allclose(M_lm.fittedvalues, M2_lm.fittedvalues)
```

```
[ ]: True
```

```
[ ]: # Extract the coefficients for month
coef_month = S2[S2.index.str.contains('mnth']]['coef']
coef_month
```


```
[ ]: mnth[Jan]      -46.0871
      mnth[Feb]    -39.2419
      mnth[March] -29.5357
      mnth[April]  -4.6622
      mnth[May]    26.4700
      mnth[June]   21.7317
      mnth[July]   -0.7626
      mnth[Aug]    7.1560
      mnth[Sept]   20.5912
      mnth[Oct]    29.7472
      mnth[Nov]    14.2229
      Name: coef, dtype: float64
```

```
[ ]: # Append Dec as the negative of the sum of other months
months = Bike['mnth'].dtype.categories
coef_month = pd.concat([
```

```
coef_month, coef_month , pd.Series([-coef_month.sum()])
])
```

```
[ ]: try:
    fig_month, ax_month = subplots(figsize=(8,8))
    x_month = np.arange(coef_month.shape[0])
    ax_month.plot(x_month, coef_month, marker='o', ms=10)
    ax_month.set_xticks(x_month)
    ax_month.set_xticklabels([l[5] for l in coef_month.index], fontsize
=20)
    ax_month.set_xlabel('Month', fontsize=20)
    ax_month.set_ylabel('Coefficient', fontsize=20);
except Exception as e:
    print(e)
```

'int' object is not subscriptable





```
[ ]: coef_hr = S2[S2.index.str.contains('hr')]['coef']
coef_hr = coef_hr.reindex(['hr[{0}]'.format(h) for h in range(23)])
coef_hr = pd.concat([coef_hr, pd.Series([-coef_hr.sum()], index=['hr[23]'])])
```

```
[ ]: fig_hr, ax_hr = subplots(figsize=(8,8))
x_hr = np.arange(coef_hr.shape[0])
ax_hr.plot(x_hr, coef_hr, marker='o', ms=10)
ax_hr.set_xticks(x_hr[::2])
ax_hr.set_xticklabels(range(24)[::2], fontsize=20)
ax_hr.set_xlabel('Hour', fontsize=20)
ax_hr.set_ylabel('Coefficient', fontsize=20);
```

**Poisson Regression** Instead of a linear regression we use poisson regression, which is more suitable for this data.



```
[ ]: M_pois = sm.GLM(Y, X2, family=sm.families.Poisson()).fit()
```

```
[ ]: S_pois = summarize(M_pois)
coef_month = S_pois[S_pois.index.str.contains('mnth')]['coef']
coef_month = pd.concat([coef_month,
pd.Series([-coef_month.sum()], index=['mnth[Dec]'])])
coef_hr = S_pois[S_pois.index.str.contains('hr')]['coef']
coef_hr = pd.concat([coef_hr, pd.Series([-coef_hr.sum()],
index=['hr[23]'])])
```

The following plots are of the coefficients of mnth and hr.

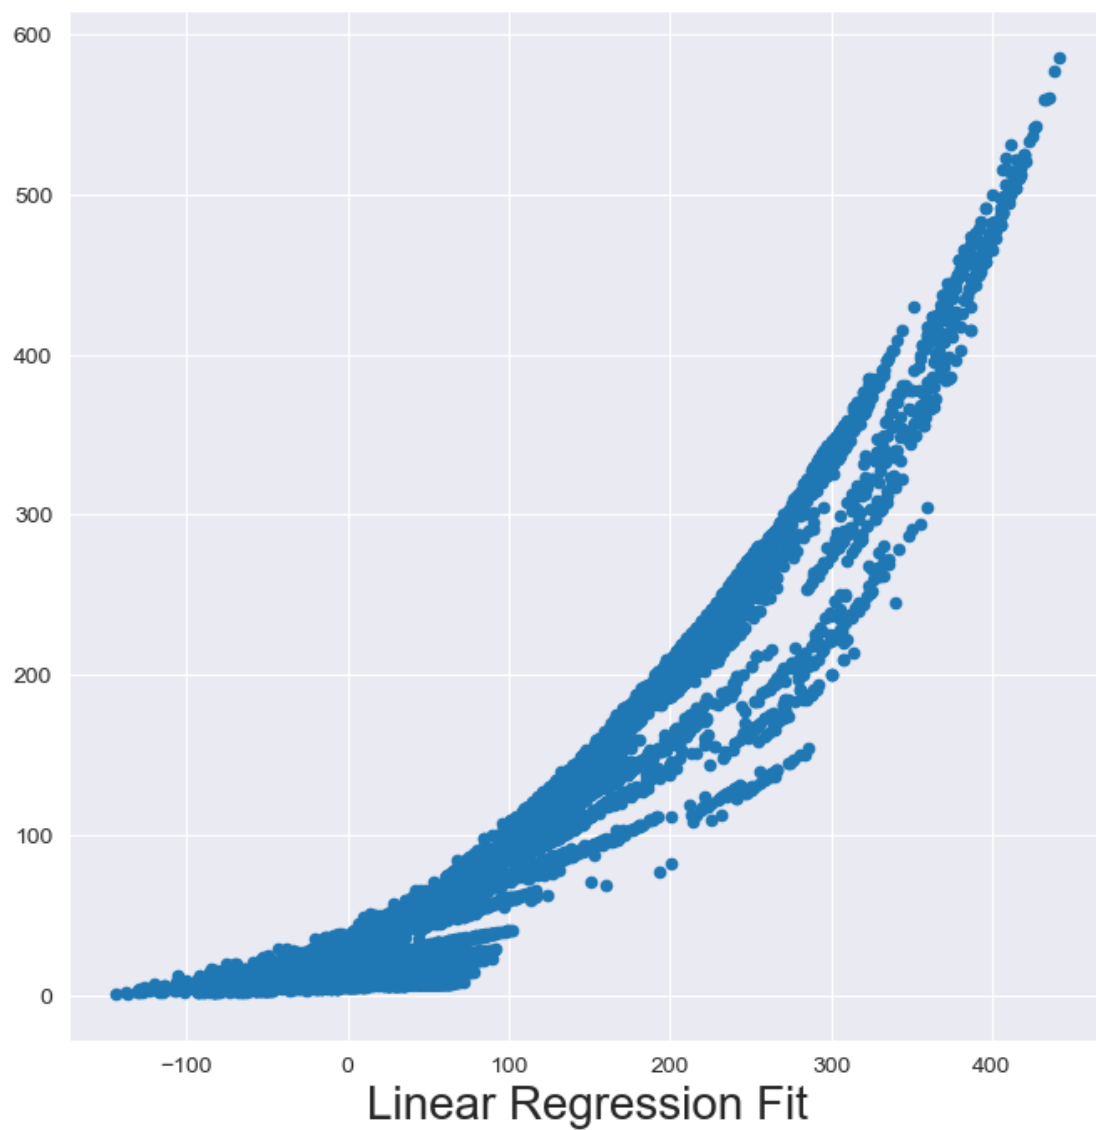
```
[ ]: fig_pois, (ax_month, ax_hr) = subplots(1, 2, figsize=(16,8))
x_month = np.arange(coef_month.shape[0])
x_hr = np.arange(coef_hr.shape[0])
ax_month.plot(x_month, coef_month, marker='o', ms=10)
ax_month.set_xticks(x_month)
ax_month.set_xticklabels([l[5] for l in coef_month.index], fontsize=20)
ax_month.set_xlabel('Month', fontsize=20)
ax_month.set_ylabel('Coefficient', fontsize=20)
ax_hr.plot(x_hr, coef_hr, marker='o', ms=10)
ax_hr.set_xticklabels(range(24)[::2], fontsize=20)
ax_hr.set_xlabel('Hour', fontsize=20)
ax_hr.set_ylabel('Coefficient', fontsize=20);
```

```
/var/folders/gf/bt25hkv172n_bttx0h72_6340000gn/T/ipykernel_12419/824882860.py:11
: UserWarning: FixedFormatter should only be used together with FixedLocator
ax_hr.set_xticklabels(range(24)[::2], fontsize=20)
```



```
[ ]: fig, ax = subplots(figsize=(8, 8))
      ax.scatter(M2_lm.fittedvalues , M_pois.fittedvalues ,
                 s=20)
      ax.set_xlabel('Linear Regression Fit', fontsize=20)
```

```
[ ]: Text(0.5, 0, 'Linear Regression Fit')
```



This shows that the prediction from the poisson regressions are correlated predictions from the linear model.