analysis

December 13, 2023

Spotify Song Data Analysis   Thursday, December 14th Jack Krebsbach & Eli Edwards

Import all libraries needed for analysis   We will use various models to predict song popularity and classify which playlist genre a song came from.

Column Name

Data Type

Description

playlist_genre

character

Playlist genre

playlist_subgenre

character

Playlist subgenre

danceability

double

Danceability describes how suitable a track is for dancing based on a combination of musical elements including tempo, rhythm stability, beat strength, and overall regularity. A value of 0.0 is least danceable and 1.0 is most danceable.

energy

double

Energy is a measure from 0.0 to 1.0 and represents a perceptual measure of intensity and activity. Typically, energetic tracks feel fast, loud, and noisy. For example, death metal has high energy, while a Bach prelude scores low on the scale. Perceptual features contributing to this attribute include dynamic range, perceived loudness, timbre, onset rate, and general entropy.

key

double

The estimated overall key of the track. Integers map to pitches using standard Pitch Class notation. E.g., 0 = C, 1 = C /D , 2 = D, and so on. If no key was detected, the value is -1.

loudness

double

The overall loudness of a track in decibels (dB). Loudness values are averaged across the entire track and are useful for comparing relative loudness of tracks. Loudness is the quality of a sound that is the primary psychological correlate of physical strength (amplitude). Values typically range between -60 and 0 dB.

mode

double

Mode indicates the modality (major or minor) of a track, the type of scale from which its melodic content is derived. Major is represented by 1 and minor is 0.

speechiness

double

Speechiness detects the presence of spoken words in a track. The more exclusively speech-like the recording (e.g., talk show, audio book, poetry), the closer to 1.0 the attribute value. Values above 0.66 describe tracks that are probably made entirely of spoken words. Values between 0.33 and 0.66 describe tracks that may contain both music and speech, either in sections or layered, including such cases as rap music. Values below 0.33 most likely represent music and other non-speech-like tracks.

acousticness

double

A confidence measure from 0.0 to 1.0 of whether the track is acoustic. 1.0 represents high confidence the track is acoustic.

instrumentalness

double

Predicts whether a track contains no vocals. "Ooh" and "aah" sounds are treated as instrumental in this context. Rap or spoken word tracks are clearly "vocal". The closer the instrumentalness value is to 1.0, the greater likelihood the track contains no vocal content. Values above 0.5 are intended to represent instrumental tracks, but confidence is higher as the value approaches 1.0.

liveness

double

Detects the presence of an audience in

Imports Generally we will use `sklearn` for model building. We will use `pandas` to read in the data.

```
# Imports for analysis
import pandas as pd
import itertools
import xgboost as xgb
```

```python
import numpy as np
import seaborn as sns
from sklearn.decomposition import PCA
from sklearn.model_selection import GridSearchCV
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
from sklearn.preprocessing import LabelEncoder
from ISLP import confusion_table
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import (RandomForestRegressor as RF,␣
 ↪RandomForestClassifier as RFC, GradientBoostingRegressor as GBR)
import sklearn.model_selection as skm
from sklearn.preprocessing import StandardScaler
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from sklearn.tree import ( DecisionTreeRegressor as DTR, plot_tree)
from sklearn.metrics import (accuracy_score,r2_score)
```

Import Data

```python
# Read in the CSV
df = pd.read_csv("./clean_data/spotify_songs.csv")

# Transform categorical data
df['playlist_genre'] = df['playlist_genre'].astype('category')
df['key'] = df['key'].astype('category')
df['mode'] = df['mode'].astype('category')

scaler = StandardScaler()
# Standardize numerical data
numeric_cols = df.select_dtypes(include=['number'])
df[numeric_cols.columns] = numeric_cols.astype('float64')
df[numeric_cols.columns] = scaler.fit_transform(numeric_cols)
```

```python
# Look at first few columns
df.head()
```

```
                   track_id                                   track_name  \
0  6f807x0ima9a1j3VPbc7VN   I Don't Care (with Justin Bieber) - Loud Luxur…
1  0r7CVbZTWZgbTCYdfa2P31                    Memories - Dillon Francis Remix
2  1z1Hg7Vb0AhHDiEmnDE79l                     All the Time - Don Diablo Remix
3  75FpbthrwQmzHlBJLuGdC7                     Call You Mine - Keanu Silva Remix
4  1e8PAfcKUYoKkxPhrHqw4x         Someone You Loved - Future Humans Remix

       track_artist  track_popularity           track_album_id  \
0         Ed Sheeran          0.941531   2oCs0DGTsRO98Gh5ZSl2Cx
```

3

```
1        Maroon 5           0.981557  63rPSO264uRjW1X5E6cWv6
2      Zara Larsson         1.101635  1HoSmj2eLcsrR0vE9gThr4
3  The Chainsmokers         0.701374  1nqYsOef1yKKuGOVchbsk6
4    Lewis Capaldi          1.061609  7m7vv9wlQ4i0LFuJiE2zsQ

                                  track_album_name track_album_release_date  \
0  I Don't Care (with Justin Bieber) [Loud Luxury…                2019-06-14
1                  Memories (Dillon Francis Remix)                2019-12-13
2                  All the Time (Don Diablo Remix)                2019-07-05
3                     Call You Mine - The Remixes                 2019-07-19
4             Someone You Loved (Future Humans Remix)             2019-03-05

  playlist_name           playlist_id playlist_genre  … key  loudness  \
0    Pop Remix  37i9dQZF1DXcZDD7cfEKhW            pop  …   6  1.367123
1    Pop Remix  37i9dQZF1DXcZDD7cfEKhW            pop  …  11  0.585766
2    Pop Remix  37i9dQZF1DXcZDD7cfEKhW            pop  …   1  1.100090
3    Pop Remix  37i9dQZF1DXcZDD7cfEKhW            pop  …   7  0.984309
4    Pop Remix  37i9dQZF1DXcZDD7cfEKhW            pop  …   1  0.685151

   mode speechiness acousticness instrumentalness  liveness  valence  \
0     1   -0.481362    -0.333898        -0.377953 -0.809230  0.031908
1     1   -0.688642    -0.468670        -0.359177  1.081061  0.782522
2     0   -0.324422    -0.436799        -0.377849 -0.519562  0.439384
3     1   -0.050024    -0.667642        -0.377911  0.089582 -1.001795
4     1   -0.702460    -0.432701        -0.377953 -0.692585  0.919777

      tempo  duration_ms
0  0.042927    -0.518874
1 -0.777198    -1.056268
2  0.116227    -0.822017
3  0.039953    -0.947750
4  0.115037    -0.614172

[5 rows x 23 columns]
```
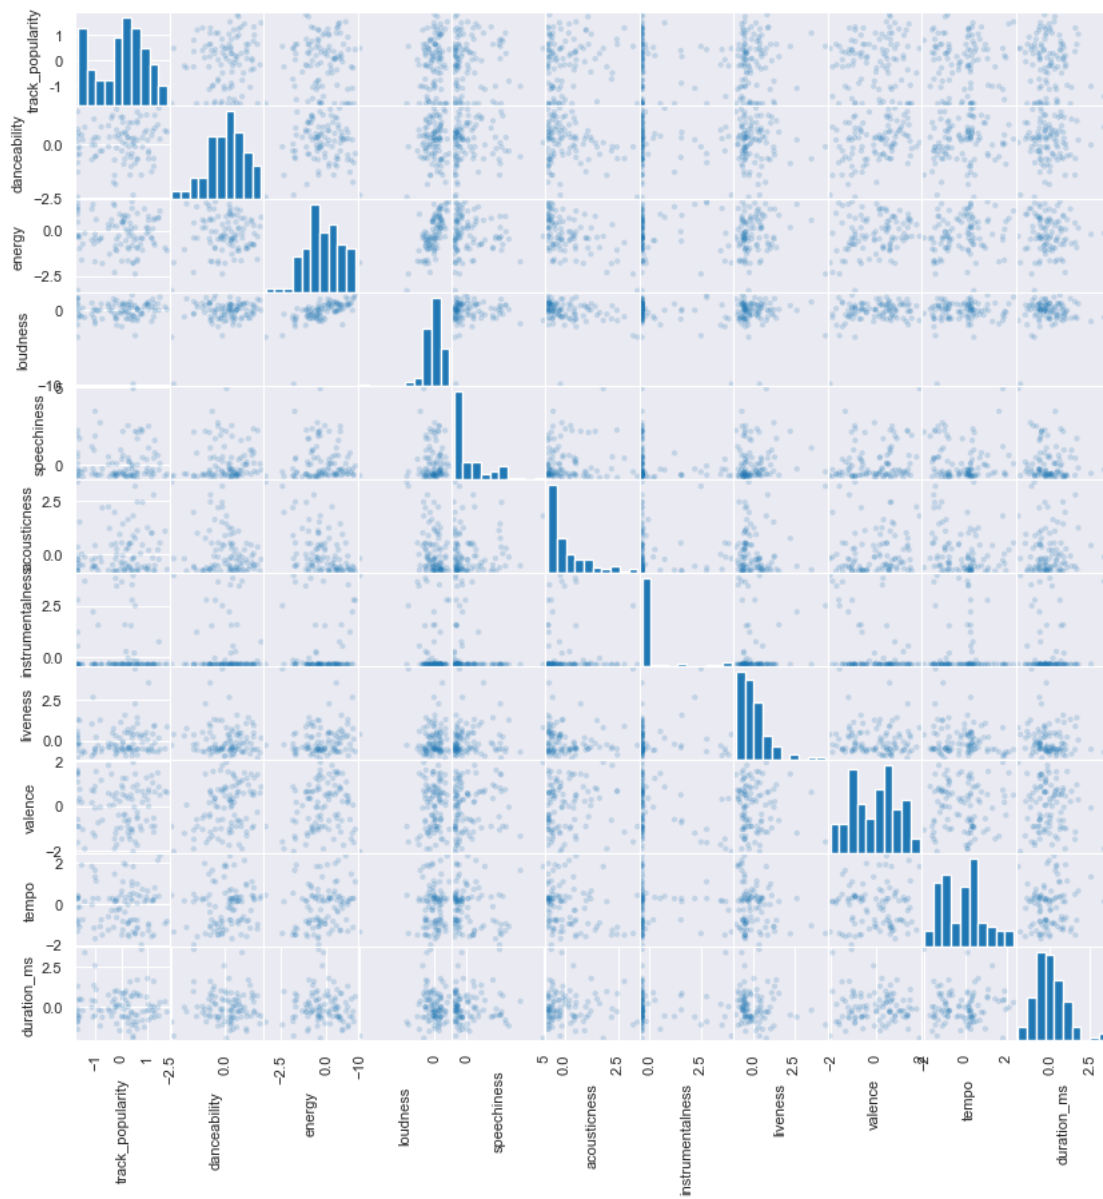
0.0.1 Scatter plot of data

```
[ ]: sampled_df = df[numeric_cols.columns].sample(n=100, random_state=42)

     # Create a scatter matrix from the sampled data
     axes = pd.plotting.scatter_matrix(sampled_df, alpha=0.2, figsize=(10, 10))

     # Adjust label properties
     for ax in axes.flatten():
         ax.xaxis.label.set_rotation(90)
         ax.xaxis.label.set_fontsize(8)
         ax.yaxis.label.set_fontsize(8)
```
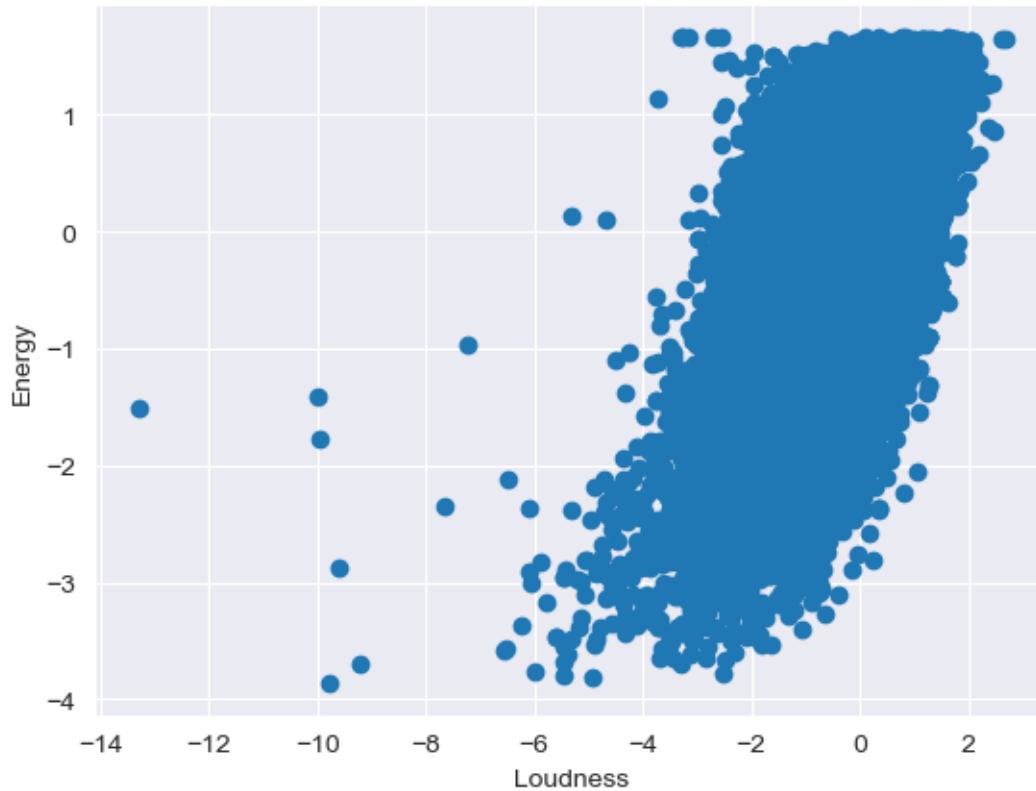
```
# Display the plot
plt.show()
```



```
plt.scatter(df.loudness, df.energy)
plt.xlabel('Loudness')
plt.ylabel('Energy')
plt.title = 'Loudness vs Energy'
plt.show()
```

Regression Analysis   We start out by splitting our data in training and test sets

```
(train, test) = skm.train_test_split(df,
                                     test_size=0.3,
                                     random_state=0)
```

```
# Predictors
cols = [ 'danceability',
        'energy',
        'key',
        'loudness',
        'mode',
        'speechiness',
        'acousticness',
        'instrumentalness',
        'liveness',
        'valence',
        'tempo',
        'duration_ms']
```

Simple Linear Regression We fit a linear model using our training data. While the coefficients appear to be significant (aside from key and mode), the R squared value is 0.072, indicating that our variables don't show much correlation. When testing the model on our test set, its mean squared error values close to 1 - which implies poor prediction because our data was scaled by PCA. Looking at a graph of the residuals vs fitted values, the data appears to be resting on a slanted slope. This is irregular and likely means that our selected variables aren't explaining enough.

```
[ ]: # Training
     x_train = train[cols]
     popularity_train = train.track_popularity
```

```
[ ]: # Testing
     x_test = test[cols]
     popularity_test = test.track_popularity
```

```
[ ]: # Linear Model
     reg = sm.OLS(popularity_train, x_train)
     results = reg.fit()
     results.summary()
```

[ ]:

Dep. Variable:	track_popularity	R-squared (uncentered):	0.072
Model:	OLS	Adj. R-squared (uncentered):	0.072
Method:	Least Squares	F-statistic:	148.5
Date:	Wed, 13 Dec 2023	Prob (F-statistic):	0.00
Time:	22:34:49	Log-Likelihood:	-31785.
No. Observations:	22983	AIC:	6.359e+04
Df Residuals:	22971	BIC:	6.369e+04
Df Model:	12		
Covariance Type:	nonrobust		

	coef	std err	t	P> |t|	[0.025	0.975]
danceability	0.0309	0.007	4.232	0.000	0.017	0.045
energy	-0.2025	0.010	-19.534	0.000	-0.223	-0.182
key	-0.0013	0.001	-1.101	0.271	-0.004	0.001
loudness	0.1804	0.009	19.577	0.000	0.162	0.198
mode	0.0167	0.010	1.626	0.104	-0.003	0.037
speechiness	-0.0337	0.007	-5.128	0.000	-0.047	-0.021
acousticness	0.0411	0.008	5.336	0.000	0.026	0.056
instrumentalness	-0.1061	0.007	-15.690	0.000	-0.119	-0.093
liveness	-0.0235	0.007	-3.601	0.000	-0.036	-0.011
valence	0.0212	0.007	2.969	0.003	0.007	0.035
tempo	0.0273	0.007	4.150	0.000	0.014	0.040
duration_ms	-0.1135	0.007	-17.434	0.000	-0.126	-0.101

| Omnibus:	2726.957	Durbin-Watson:	2.006
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1002.332
Skew:	-0.288	Prob(JB):	2.22e-218
Kurtosis:	2.155	Cond. No.	13.7

Notes:

[1] R^2 is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```python
# Fitting our test data
predict = results.get_prediction(x_test)
predicted = predict.predicted_mean
np.mean((popularity_test - predicted)**2)
```
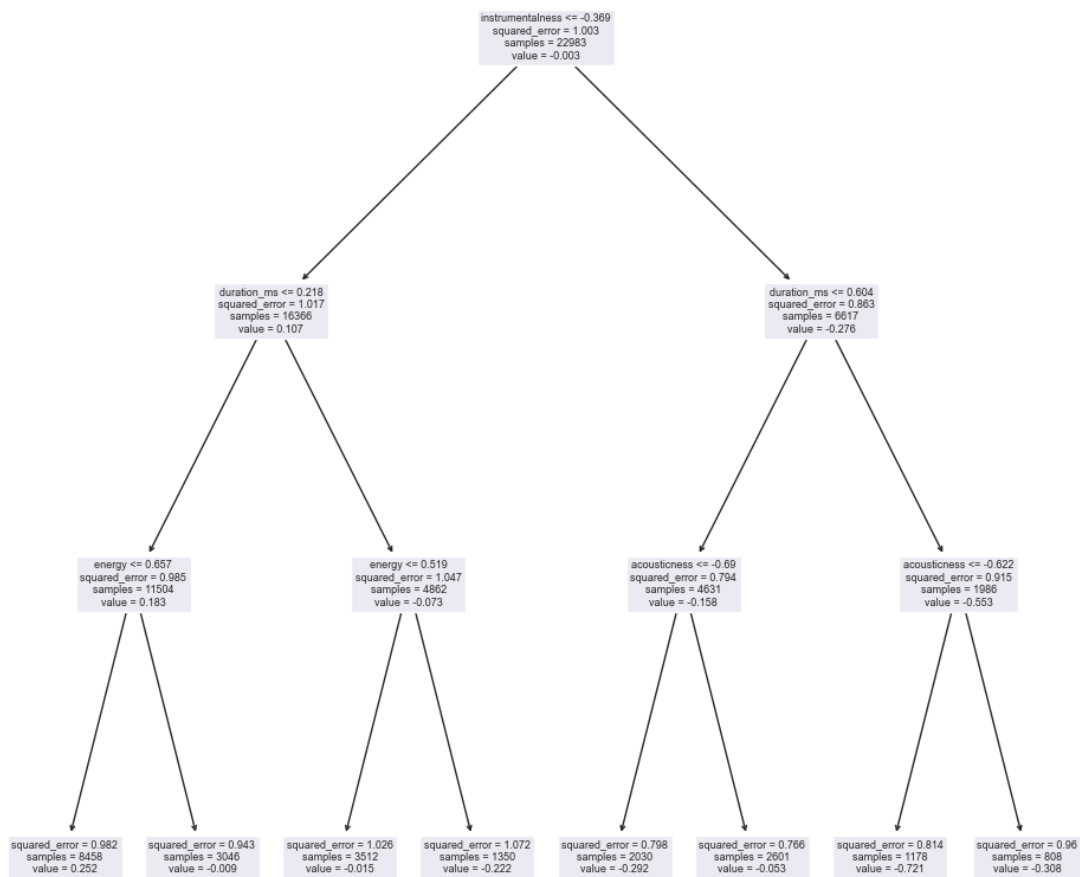
```
0.9222819329563361
```

```python
# Plotting residuals vs fitted values
ax = subplots(figsize=(8,8))[1]
ax.scatter(results.fittedvalues , results.resid)
ax.set_xlabel('Fitted value')
ax.set_ylabel('Residual')
ax.axhline(0, c='k', ls='--');
```
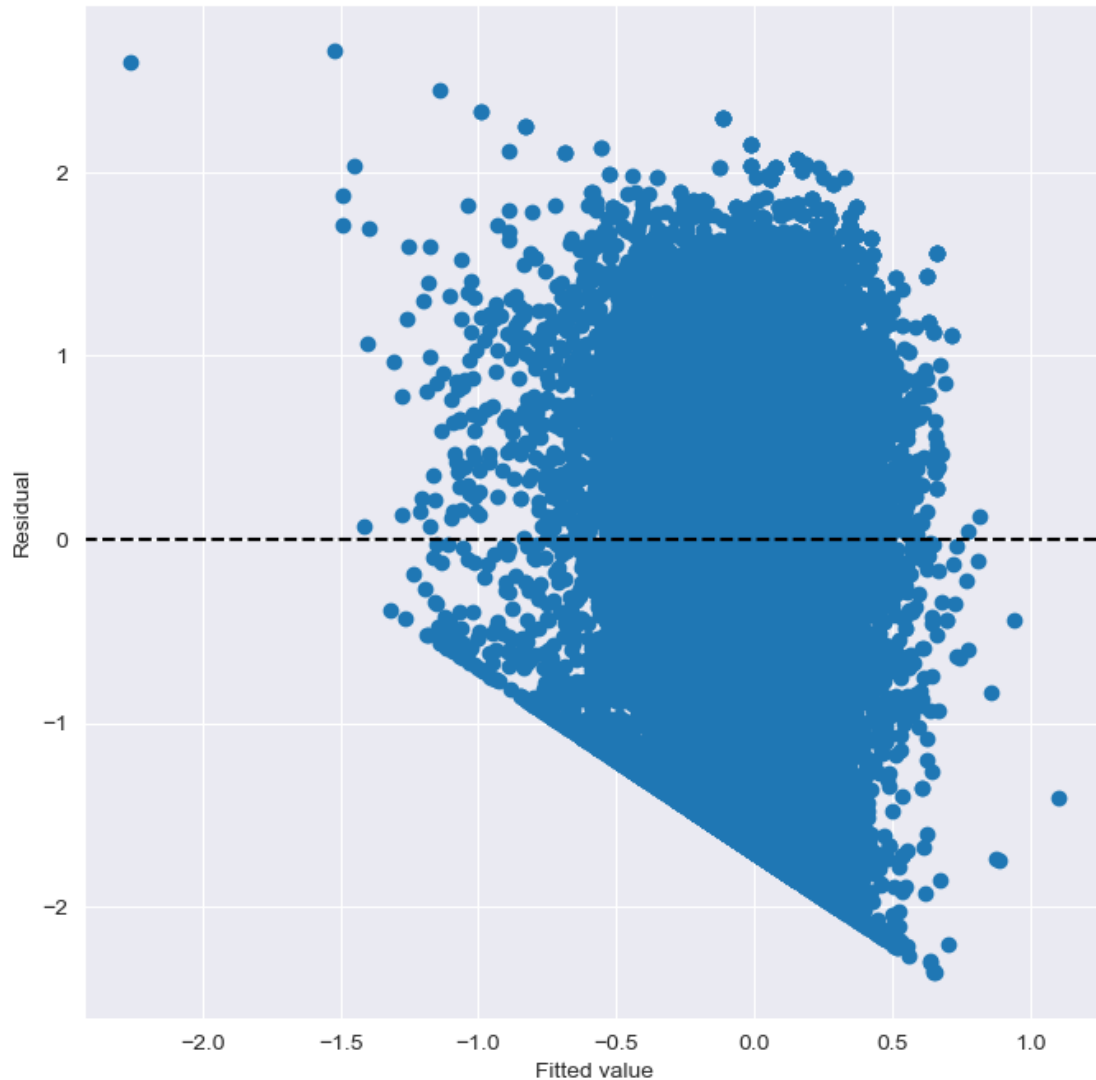
Regression Tree Analysis   We then fit our data to a regression tree model. We pruned the tree to find the optimal model, which unfortunately did not do a very good job at predicting the data. We fit the data to the test set, and got an even lower R squared value than before. The mean squared error was again close to 1. Getting the importance values, instrumentalness and duration seem to have the most impact on our results, even though they are not very significant.

```
# Creating initial tree
reg = DTR(max_depth=3)
reg.fit(x_train, popularity_train)
ax = subplots(figsize=(12,12))[1]
plot_tree(reg,
          feature_names=reg.feature_names_in_,
          ax=ax);
```

9

instrumentalness <= -0.369
squared_error = 1.003
samples = 22983
value = -0.003

duration_ms <= 0.218
squared_error = 1.017
samples = 16366
value = 0.107

duration_ms <= 0.604
squared_error = 0.863
samples = 6617
value = -0.276

energy <= 0.657
squared_error = 0.985
samples = 11504
value = 0.183

energy <= 0.519
squared_error = 1.047
samples = 4862
value = -0.073

acousticness <= -0.69
squared_error = 0.794
samples = 4631
value = -0.158

acousticness <= -0.622
squared_error = 0.915
samples = 1986
value = -0.553

squared_error = 0.982
samples = 8458
value = 0.252

squared_error = 0.943
samples = 3046
value = -0.009

squared_error = 1.026
samples = 3512
value = -0.015

squared_error = 1.072
samples = 1350
value = -0.222

squared_error = 0.798
samples = 2030
value = -0.292

squared_error = 0.766
samples = 2601
value = -0.053

squared_error = 0.814
samples = 1178
value = -0.721

squared_error = 0.96
samples = 808
value = -0.308

```python
# 'Pruning' the tree
ccp_path = reg.cost_complexity_pruning_path(x_train, popularity_train)
kfold = skm.KFold(5,
                  shuffle=True,
                  random_state=10)
grid = skm.GridSearchCV(reg,
                        {'ccp_alpha': ccp_path.ccp_alphas},
                        refit=True,
                        cv=kfold,
                        scoring='neg_mean_squared_error')
G = grid.fit(x_train, popularity_train)
```

```
[ ]: # Fitting our test data
     best_ = grid.best_estimator_
     predicted = best_.predict(x_test)

     # MSE
     print(np.mean((popularity_test - best_.predict(x_test))**2))

     # R squared
     r2_score(popularity_test, predicted)
```

0.9443524899017082

```
[ ]: 0.049260738158203754
```

```
[ ]: # Importance values of each predictor
     print(best_.feature_importances_)
     best_.feature_names_in_
```

```
[0.          0.13169085 0.          0.          0.          0.
 0.09989581 0.46945227 0.          0.          0.          0.29896106]
```

```
[ ]: array(['danceability', 'energy', 'key', 'loudness', 'mode', 'speechiness',
            'acousticness', 'instrumentalness', 'liveness', 'valence', 'tempo',
            'duration_ms'], dtype=object)
```

Classification Analysis Now we will explore classifying the playlist genre these songs came from.

Overall, the XGBoost classifier had the highest accuracy on the test set. We use k-vfold with 5 folds over a grid of parameters.

XGBoost model - Max Depth: 5 - Number of estimators: 300 - Learning rate: 0.1 - Test set accuracy: 57.4%

Speechiness was the highest importance

0.0.2 Plot first two principal components of data

To get to know the data first we plot the first two principal components.

Generally the data are scattered across classes. EDM seems to be the most clustered class when plotted on the first two principal components.

```
[ ]: # Get the first two principal components
     pca = PCA(n_components=2)
     principal_components = pca.fit_transform(numeric_cols)
     pca_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])
     pca_df['playlist_genre'] = df['playlist_genre'].values  # Ensure this column␣
      ↪aligns with your PCA data
```

```python
# Sample 150 points from each class for plotting
plot_df = pd.DataFrame()
for genre in pca_df['playlist_genre'].unique():
    sampled_genre_df = pca_df[pca_df['playlist_genre'] == genre].sample(n=150,␣
 ↪random_state=42)
    plot_df = pd.concat([plot_df, sampled_genre_df], axis=0)

# Plot the results
plt.figure(figsize=(10, 8))
sns.scatterplot(x='PC1', y='PC2', hue='playlist_genre', data=plot_df)
plt.title('PCA of Dataset (Sampled for Visualization)')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend()
plt.show()
```

```
---------------------------------------------------------------------------
NameError                                 Traceback (most recent call last)
Cell In[1], line 2
      1 # Get the first two principal components
----> 2 pca = PCA(n_components=2)
      3 principal_components = pca.fit_transform(numeric_cols)
      4 pca_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2']

NameError: name 'PCA' is not defined
```

Create train/test split   Because we have so many observations we can afford to only use 10% of the data on testing. We stratify the data be playlist genre to ensure that all classes are included an equal amount in the training and testing data.

```python
(train, test) = skm.train_test_split(df, test_size=0.10, random_state=0,␣
 ↪stratify=df.playlist_genre)
```

0.0.3   Make model matrix for classification

```python
# Predictors
cols = ['track_popularity',
        'danceability',
        'energy',
        'key',
        'loudness',
        'mode',
        'speechiness',
        'acousticness',
        'instrumentalness',
        'liveness',
```

```
            'valence',
            'tempo',
            'duration_ms']
```

```
[ ]:  # Training
      x_train = train[cols]
      genre_train = train.playlist_genre
```

```
[ ]:  # Testing
      x_test = test[cols]
      genre_test = test.playlist_genre
```

0.0.4 K-Nearest Neighbors

For a baseline, we try and fit a K-Nearest-Neighbors model.

```
[ ]:  # Use approximately the square root of number of observations
      #k = int(np.floor(np.sqrt(x_train.shape[0])))
      k = 20
      print(f'Using {k} neighbors')
      # Initialize classifier
      knn = KNeighborsClassifier(n_neighbors=k)
      # Fit the model
      knn.fit(x_train, genre_train)
      # Predict on the testing set
      knn_pred = knn.predict(x_test)
      # Get confusion table
      confusion_table(knn_pred, genre_test)
```

```
      Using 20 neighbors
```

```
[ ]:  Truth      edm  latin  pop  r&b  rap  rock
      Predicted
      edm        270    148  161  132  150   117
      latin       67     76   87   92   53    55
      pop         90     99  125   69   90    74
      r&b         58     60   65  104  100    90
      rap         71     73   54   79  124    57
      rock        48     60   59   67   58   102
```

```
[ ]:  # Get confusion matrix
      np.mean(knn_pred == genre_test)
      accuracy_score(genre_test, knn_pred)
```

```
[ ]:  0.24390986601705236
```

Tune KNN Model   We will tune the KNN to find the optimal K from 1 to $\lfloor \sqrt{n} \rfloor + 1$ where n is the number of observations in the test set.

```python
# Clear plot
plt.clf()
# Range of k to try
k_range = range(1, int(np.floor(np.sqrt(x_train.shape[0]))) + 1)

# Lists to store accuracies
train_error = []
test_error = []
# Loop over different values of k
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, genre_train)

    # Training accuracy
    train_pred = knn.predict(x_train)
    train_error.append(1- accuracy_score(genre_train, train_pred))

    # Testing accuracy
    test_pred = knn.predict(x_test)
    test_error.append(1- accuracy_score(genre_test, test_pred))

# Plotting
plt.figure(figsize=(10, 6))
plt.plot(k_range, train_error, label='Training Error')
plt.plot(k_range, test_error, label='Testing Error')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Error')
plt.title('KNN Training and Testing Error')
plt.legend()
plt.show()
```

```
<Figure size 640x480 with 0 Axes>
```

KNN Training and Testing Error

It looks like the best that the KNN classifier can do is around 25%.

LDA Classifier Now we fit a linear discriminant classifier to classify the song genres. There are no hyper parameters to tune.

```
# Instantiate the LDA model
lda = LinearDiscriminantAnalysis()
# Fit the model
lda.fit(x_train, genre_train)
# Predict on new data
lda_pred = lda.predict(x_test)
# Confusion matrix
cm = confusion_matrix(genre_test, lda_pred)
```

```
# Get genre labels
genre_labels = genre_test.cat.categories.tolist()

# Coerce into data frame
cm_df = pd.DataFrame(cm, index=genre_labels, columns=genre_labels)
# Plot the DataFrame using matplotlib
fig, ax = plt.subplots(figsize=(5, 5))  # Set figure size
ax.axis('off')
```

```
tbl = ax.table(cellText=cm_df.values, colLabels=cm_df.columns, rowLabels=cm_df.
  ↪index, loc='center', cellLoc='center')
tbl.auto_set_font_size(False)
tbl.set_fontsize(14)
tbl.scale(1.5, 1.5)

# Add labels for True Label and Predicted Label in the appropriate position
ax.text(-0.40, 0.5, 'Truth', va='center', ha='center', rotation='vertical',␣
  ↪size=16, transform=ax.transAxes)
ax.text(0, 0.75, 'Predicted',size=16)
print(f'Class Accuracies: {np.diag(cm / np.sum(cm, axis=0))}')
#print(f'Normalized Accuracies: {cm / np.sum(cm, axis=0)}')
```

Class Accuracies: [0.55156951 0.37452471 0.36051502 0.44329897 0.50592217
0.54844607]

Predicted

	edm	latin	pop	r&b	rap	rock
edm	369	59	89	19	39	29
latin	59	197	65	59	94	42
pop	81	96	168	71	39	96
r&b	29	73	56	215	117	53
rap	78	73	37	61	299	27
rock	53	28	51	60	3	300

(Truth labels the rows)

```
[ ]: # Testing Accuracy
accuracy_score(genre_test, lda_pred)
```

[ ]: 0.4713763702801462

Plot the decision boundary for PCA Data   Using the first two principal components we fit
the data and plot the decision boundaries.

```python
# First fit another scaled PCA and predict probabilities
scaler = StandardScaler()
# Fit the scaler to the PCA components and transform the data
# Might have already been done, so we catch the error
try:
    pca_df = pca_df.drop('playlist_genre', axis=1)
except Exception as e:
    print()

pca_df[['PC1', 'PC2']] = scaler.fit_transform(pca_df[['PC1', 'PC2']])

# Instantiate the LDA model for PCA version
lda_pca = LinearDiscriminantAnalysis()

# Fit the model
lda_pca.fit(pca_df[['PC1', 'PC2']], df.playlist_genre)

x_min, x_max = pca_df['PC1'].min() -1 , pca_df['PC1'].max() +1
y_min, y_max = pca_df['PC2'].min() -1 , pca_df['PC2'].max() +1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01), np.arange(y_min, y_max, 0.
 ↪01))

# Create mesh
mesh_points = pd.DataFrame(np.c_[xx.ravel(), yy.ravel()], columns=['PC1','PC2']
 ↪)
# Predict the class using LDA for each point on the mesh
Z = lda_pca.predict_proba(mesh_points)
```

```python
# Plot the PCA-transformed points
plt.figure(figsize=(10, 8))

num_classes = 6
colors = plt.cm.viridis(np.linspace(0, 1, num_classes))

for i in range(num_classes):
    max_prob = np.max(Z, axis=1) == Z[:, i]
    contour = plt.contour(xx, yy, max_prob.reshape(xx.shape),
 ↪colors=[colors[i]])

sns.scatterplot(x='PC1', y='PC2', hue=df.playlist_genre, data=pca_df)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('LDA Decision Boundary in PCA Dimensions')
plt.show()
```

LDA Decision Boundary in PCA Dimensions

Random Forest   We will use Random Forest to classify the songs. With KV-Fold validation we will tune the number of estimators, the max_features used at each split, the max depth of each tree, and the minimum samples in each leaf. We will use 5 folds over the grid.

```
[ ]: # Instantiate classifier instance
     rfc = RFC(random_state=0, n_estimators=500,  min_samples_leaf=10)
     # Fit the data
     rfc.fit(x_train, genre_train)
     # Predict on new data
     rfc_pred = rfc.predict(x_test)
     # Confusion matrix
     cm = confusion_matrix(genre_test, rfc_pred)
     cm
```

```
[ ]: # Get the accuracy
     np.mean(rfc_pred == genre_test)
```

```
[ ]: # Extract the first tree from the forest
     first_tree = rfc.estimators_[0]

     # Plot the first tree
     plt.figure(figsize=(20,10))
     plot_tree(first_tree, filled=True, feature_names=x_train.columns,␣
       ↪class_names=True)
     plt.title("First Tree in the Random Forest")
     plt.show()
```

Tune Parameters for Random Forest

```
[ ]: # Define Grid
     # param_grid = {
     #     'n_estimators': [100, 200, 500],
     #     'max_features': ['sqrt', 'log2'],
     #     'max_depth': [4, 5, 6, 7, 8],
     #     'criterion': ['gini',],
     #     'min_samples_leaf': [1, 2, 4, 10]
     # }
     #
     param_grid = {
         'n_estimators': [100,500],
         'max_features': ['sqrt'],
         'max_depth': [4],
         'criterion': ['gini'],
         'min_samples_leaf': [2]
     }
```

```
[ ]: rfc = RFC(random_state=0)
     CV_rfc = GridSearchCV(estimator=rfc, param_grid=param_grid, cv=5)
     CV_rfc.fit(x_train, genre_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=0),
                  param_grid={'criterion': ['gini'], 'max_depth': [4],
                              'max_features': ['sqrt'], 'min_samples_leaf': [2],
                              'n_estimators': [100, 500]})
```

```
[ ]: print("Best Parameters found by GridSearchCV:")
     print(CV_rfc.best_params_)
```

```
Best Parameters found by GridSearchCV:
{'criterion': 'gini', 'max_depth': 4, 'max_features': 'sqrt',
'min_samples_leaf': 2, 'n_estimators': 500}
```

```
[ ]: # Create plots for different hyperparameter values
     param_pairs = list(itertools.combinations(param_grid.keys(), 2))
```

```
num_rows = len(param_pairs) // 3 + (len(param_pairs) % 3 > 0)
fig, axs = plt.subplots(num_rows, 3, figsize=(15, 5 * num_rows))
axs = axs.flatten()

# Loop over pairs of parameters
for idx, (param1, param2) in enumerate(param_pairs):
    results = pd.DataFrame(CV_rfc.cv_results_)

    # Aggregate the results
    grouped_results = results.groupby([f"param_{param1}", f"param_{param2}"]).
 ↪mean(numeric_only=True)

    # Create a pivot table for each pair of parameters
    pivot_table = grouped_results["mean_test_score"].unstack()

    # Plotting in the specified subplot
    sns.heatmap(pivot_table, annot=True, cmap="YlGnBu", ax=axs[idx])
    axs[idx].set_title(f"Grid Search Scores ({param1} vs {param2})")
    axs[idx].set_xlabel(param2)
    axs[idx].set_ylabel(param1)

# Adjust layout
plt.tight_layout()
plt.show()
```

```
---------------------------------------------------------------------------
ValueError                                Traceback (most recent call last)
Cell In[144], line 11
      8 num_rows = len(param_pairs) // 3 + (len(param_pairs) % 3 > 0)
     10 # Create a figure with subplots
---> 11 fig, axs = plt.subplots(num_rows, 3, figsize=(15, 5 * num_rows))
     12 axs = axs.flatten()
     14 # Loop over pairs of parameters

File /opt/homebrew/Caskroom/miniforge/base/envs/stats/lib/python3.11/
 ↪site-packages/matplotlib/pyplot.py:1502, in subplots(nrows, ncols, sharex,↵
 ↪sharey, squeeze, width_ratios, height_ratios, subplot_kw, gridspec_kw,↵
 ↪**fig_kw)
   1358 """
   1359 Create a figure and a set of subplots.
   1360
   (…)
   1499
   1500 """
   1501 fig = figure(**fig_kw)
-> 1502 axs =␣
 ↪fig.subplots(nrows=nrows, ncols=ncols, sharex=sharex, sharey=sharey,
```

```
       1503                    squeeze=squeeze, subplot_kw=subplot_kw,
       1504                    gridspec_kw=gridspec_kw, height_ratios=height_ratios
       1505                    width_ratios=width_ratios)
       1506 return fig, axs

File /opt/homebrew/Caskroom/miniforge/base/envs/stats/lib/python3.11/
 ↪site-packages/matplotlib/figure.py:916, in FigureBase.subplots(self, nrows,␣
 ↪ncols, sharex, sharey, squeeze, width_ratios, height_ratios, subplot_kw,␣
 ↪gridspec_kw)
        912          raise ValueError("'width_ratios' must not be defined both as "
        913                           "parameter and as key in 'gridspec_kw'")
        914     gridspec_kw['width_ratios'] = width_ratios
  --> 916 gs = self.add_gridspec(nrows, ncols, figure=self, **gridspec_kw)
        917 axs = gs.subplots(sharex=sharex, sharey=sharey, squeeze=squeeze,
        918                   subplot_kw=subplot_kw)
        919 return axs

File /opt/homebrew/Caskroom/miniforge/base/envs/stats/lib/python3.11/
 ↪site-packages/matplotlib/figure.py:1544, in FigureBase.add_gridspec(self,␣
 ↪nrows, ncols, **kwargs)
       1505 """
       1506 Return a `.GridSpec` that has this figure as a parent.  This allows
       1507 complex layout of Axes in the figure.
       (…)
       1540
       1541 """
       1543 _ = kwargs.pop('figure', None)  # pop in case user has added this…
  -> 1544 gs = GridSpec(nrows=nrows, ncols=ncols, figure=self, **kwargs)
       1545 return gs

File /opt/homebrew/Caskroom/miniforge/base/envs/stats/lib/python3.11/
 ↪site-packages/matplotlib/gridspec.py:379, in GridSpec.__init__(self, nrows,␣
 ↪ncols, figure, left, bottom, right, top, wspace, hspace, width_ratios,␣
 ↪height_ratios)
        376 self.hspace = hspace
        377 self.figure = figure
  --> 379 super().__init__(nrows, ncols,
        380                  width_ratios=width_ratios,
        381                  height_ratios=height_ratios)

File /opt/homebrew/Caskroom/miniforge/base/envs/stats/lib/python3.11/
 ↪site-packages/matplotlib/gridspec.py:49, in GridSpecBase.__init__(self, nrows␣
 ↪ncols, height_ratios, width_ratios)
         34 """
         35 Parameters
         36 ----------
       (…)
         46     If not given, all rows will have the same height.
         47 """
```

```
      48 if not isinstance(nrows, Integral) or nrows <= 0:
---> 49     raise ValueError(
      50         f"Number of rows must be a positive integer, not {nrows!r}")
      51 if not isinstance(ncols, Integral) or ncols <= 0:
      52     raise ValueError(
      53         f"Number of columns must be a positive integer, not {ncols!r}")

ValueError: Number of rows must be a positive integer, not 0
```

```
<Figure size 1500x0 with 0 Axes>
```

```
[ ]: rfc_best = CV_rfc.best_estimator_
     rfc_pred = rfc_best.predict(x_test)
     cm = confusion_matrix(genre_test, rfc_pred)
     print(cm)
```

```
[[440  32  32  25  39  36]
 [ 66 172  34  72 130  42]
 [126  94  68 104  56 103]
 [ 31  66  25 202 157  62]
 [ 55  50  18  61 357  34]
 [ 55  19  21  71   5 324]]
```

```
[ ]: accuracy_score(genre_test, rfc_pred)
```

```
[ ]: 0.4759439707673569
```

```
[ ]: # Get feature importance for random forest
     feature_importances = rfc_best.feature_importances_

     # If x_train is a DataFrame, get the feature names
     feature_names = x_train.columns

     # Coerce to data frame
     importances_df = pd.DataFrame({'Feature': feature_names, 'Importance':␣
      ↪feature_importances})
     importances_df = importances_df.sort_values(by='Importance', ascending=False)

     # Print the top 3 features and their importances
     print("Top 3 Feature Importances:")
     print(importances_df.head(3))
```

```
Top 3 Feature Importances:
        Feature  Importance
6     speechiness    0.222504
1    danceability    0.181070
11          tempo    0.134920
```

XGBoost   To work with this classifier we first need to encode the data.

```python
# Instantiate Label Encoder
pd.options.mode.chained_assignment = None
le_genre = LabelEncoder()
le_key = LabelEncoder()

# Encode testing set
y_encoded = le_genre.fit_transform(genre_train)
x_train.loc[:,'key'] = le_key.fit_transform(x_train.key)

# Encode training set
y_test_encoded = le_genre.transform(genre_test)
x_test.loc[:,'key'] = le_key.transform(x_test.key)
```

/var/folders/gf/bt25hkv172n_bttx0h72_6340000gn/T/ipykernel_71719/1143999346.py:8
: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will
attempt to set the values inplace instead of always setting a new array. To
retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns
are non-unique, `df.isetitem(i, newvals)`
  x_train.loc[:,'key'] = le_key.fit_transform(x_train.key)
/var/folders/gf/bt25hkv172n_bttx0h72_6340000gn/T/ipykernel_71719/1143999346.py:1
2: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will
attempt to set the values inplace instead of always setting a new array. To
retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns
are non-unique, `df.isetitem(i, newvals)`
  x_test.loc[:,'key'] = le_key.transform(x_test.key)

```python
# Instantiate the classifier
xg = xgb.XGBClassifier(enable_categorical = True)
# Fit the data
xg.fit(x_train, y_encoded)
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=True, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, objective='multi:softprob', …)
```

```python
# Make prediction
xgb_pred = xg.predict(x_test)
#Confusion Matrx
```

```
cm = confusion_matrix(y_test_encoded, xgb_pred)
cm
```

[ ]: array([[427, 31, 78, 19, 30, 19],
             [32, 219, 90, 70, 81, 24],
             [95, 58, 220, 74, 35, 69],
             [18, 58, 65, 261, 100, 41],
             [26, 44, 27, 65, 388, 25],
             [16,   8, 64, 43,   5, 359]])

```
# Get accuracy
np.mean(xgb_pred == y_test_encoded)
```

[ ]: 0.5706455542021924

Now do a grid search in KV-fold cross validation To find the optimal parameters we conduct a grid search with 5 folds along various values for `max_depth`, `n_estimators`, and `learning_rate`.

```
# Define the parameter grid
param_grid = {
    'max_depth': [2,3,4, 5],
    'n_estimators': [100, 200, 300],
    'learning_rate': [0.01, 0.1]
}

# param_grid = {
#     'max_depth': [3,2],
#     'n_estimators': [100],
#     'learning_rate': [0.2]
# }
```

```
# Instantiate an XGBoost classifier
xgb_model = xgb.XGBClassifier(random_state=0,enable_categorical=True,
   ↪use_label_encoder=False, eval_metric='mlogloss')
# Create GridSearchCV object
CV_xgb = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5)
# Fit GridSearchCV
CV_xgb.fit(x_train, y_encoded)
# The best estimator after grid search
best_xgb_model = CV_xgb.best_estimator_
```

Plot Grid Search of Best Params

```
# Assuming 'param_grid' and 'CV_xgb.cv_results_' are defined
param_pairs = list(itertools.combinations(param_grid.keys(), 2))

# Calculate the number of rows needed for 3 columns
```

```
num_rows = len(param_pairs) // 3 + (len(param_pairs) % 3 > 0)

# Create a figure with subplots
fig, axs = plt.subplots(num_rows, 3, figsize=(15, 5 * num_rows))

# Flatten the array of axes for easy indexing
axs = axs.flatten()

# Loop over pairs of parameters
for idx, (param1, param2) in enumerate(param_pairs):
    results = pd.DataFrame(CV_xgb.cv_results_)

    # Aggregate the results
    grouped_results = results.groupby([f"param_{param1}", f"param_{param2}"]).
 ↪mean(numeric_only=True)

    # Create a pivot table for each pair of parameters
    pivot_table = grouped_results["mean_test_score"].unstack()

    # Plotting in the specified subplot
    sns.heatmap(pivot_table, annot=True, cmap="YlGnBu", ax=axs[idx])
    axs[idx].set_title(f"Grid Search Scores ({param1} vs {param2})")
    axs[idx].set_xlabel(param2)
    axs[idx].set_ylabel(param1)

# Adjust layout
plt.tight_layout()
plt.show()
```



```
[ ]: print("Best Parameters found by GridSearchCV:")
     print(CV_xgb.best_params_)
```

```
Best Parameters found by GridSearchCV:
{'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 300}
```

```python
# Get feature importance for xgboost
best_xgb_model = CV_xgb.best_estimator_
feature_importances = best_xgb_model.feature_importances_
feature_names = x_train.columns
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance':
  ↪feature_importances}).sort_values(by='Importance', ascending=False)

# Print the top 3 features and their importance
print("Top 3 Feature Importance:")
print(importance_df.head(3))
```

```
Top 3 Feature Importance:
          Feature  Importance
6      speechiness    0.137885
11           tempo    0.134520
1     danceability    0.123001
```

```python
xgb_best = CV_xgb.best_estimator_
xgb_pred_best = xgb_best.predict(x_test)
cm = confusion_matrix(y_test_encoded, xgb_pred_best)
print(cm)
accuracy_score(y_test_encoded, xgb_pred_best)
```

```
[[429  27  75  22  31  20]
 [ 34 221  80  66  81  34]
 [ 93  63 208  74  37  76]
 [ 21  59  58 270 104  31]
 [ 26  46  28  60 396  19]
 [ 18   6  61  39   8 363]]
```

```
0.5746041412911084
```